

REAL-WORLD DATASETS DISTRIBUTION




Figure 1: Amzn uint32




Figure 2: CompanyNet




Figure 3: Friendster




Figure 4: Wiki uint32

Real-world datasets - distribution. The x axis reports the indexes (positions) in the vector to store, while the y axis reports the corresponding values.

SYNTHETIC DATASETS DISTRIBUTION




Figure 5: Normal




Figure 6: Lognormal




Figure 7: Exponential




Figure 8: Zipf

Synthetic dataset - distribution. The x axis reports the indexes (positions) in the vector to store, while the y axis reports the corresponding values.

64-BIT DATASETS DISTRIBUTION




Figure 9: Amzn uint64




Figure 10: Facebook uint64




Figure 11: OSM Cellids uint64




Figure 12: Wiki uint64

64-bits datasets - dataset distributions. The x axis reports the indexes (positions) in the vector to store, while the y axis reports the corresponding values.

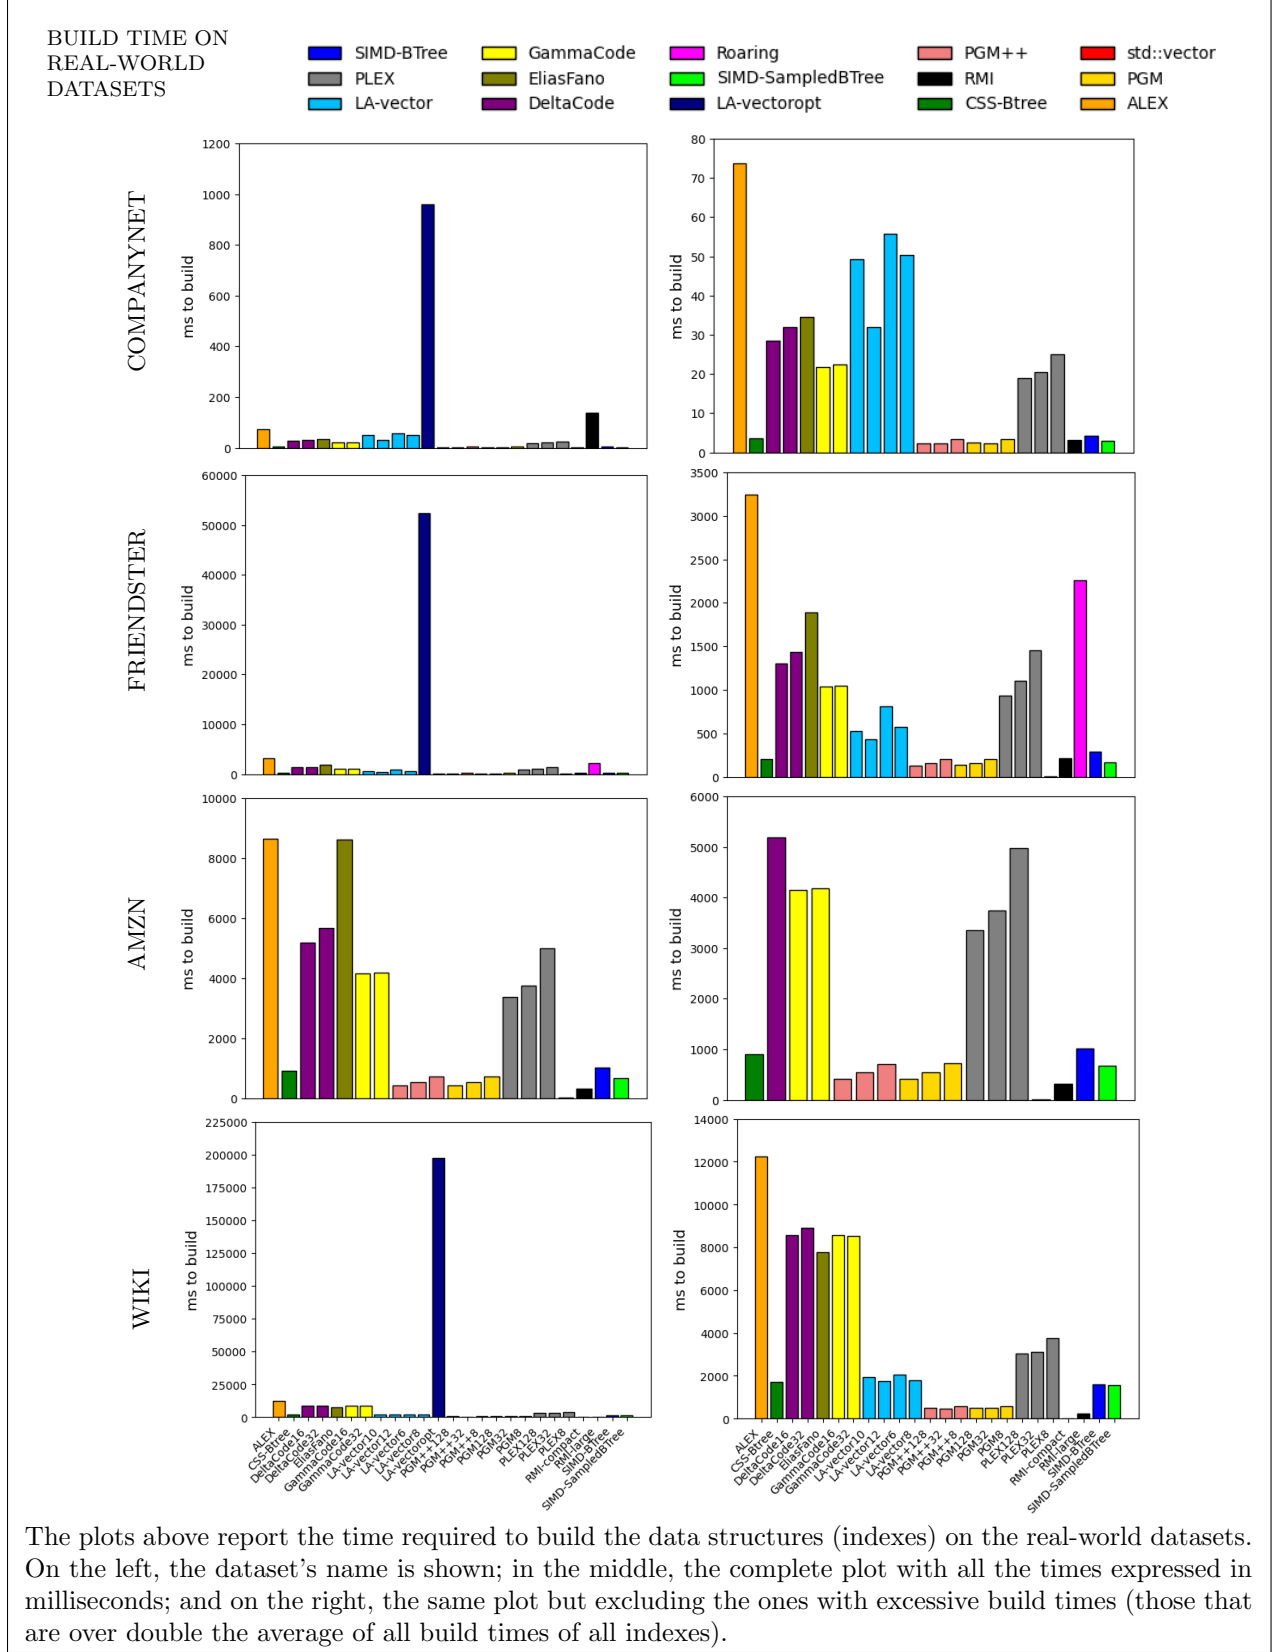


Figure 13: Build time on real-world datasets

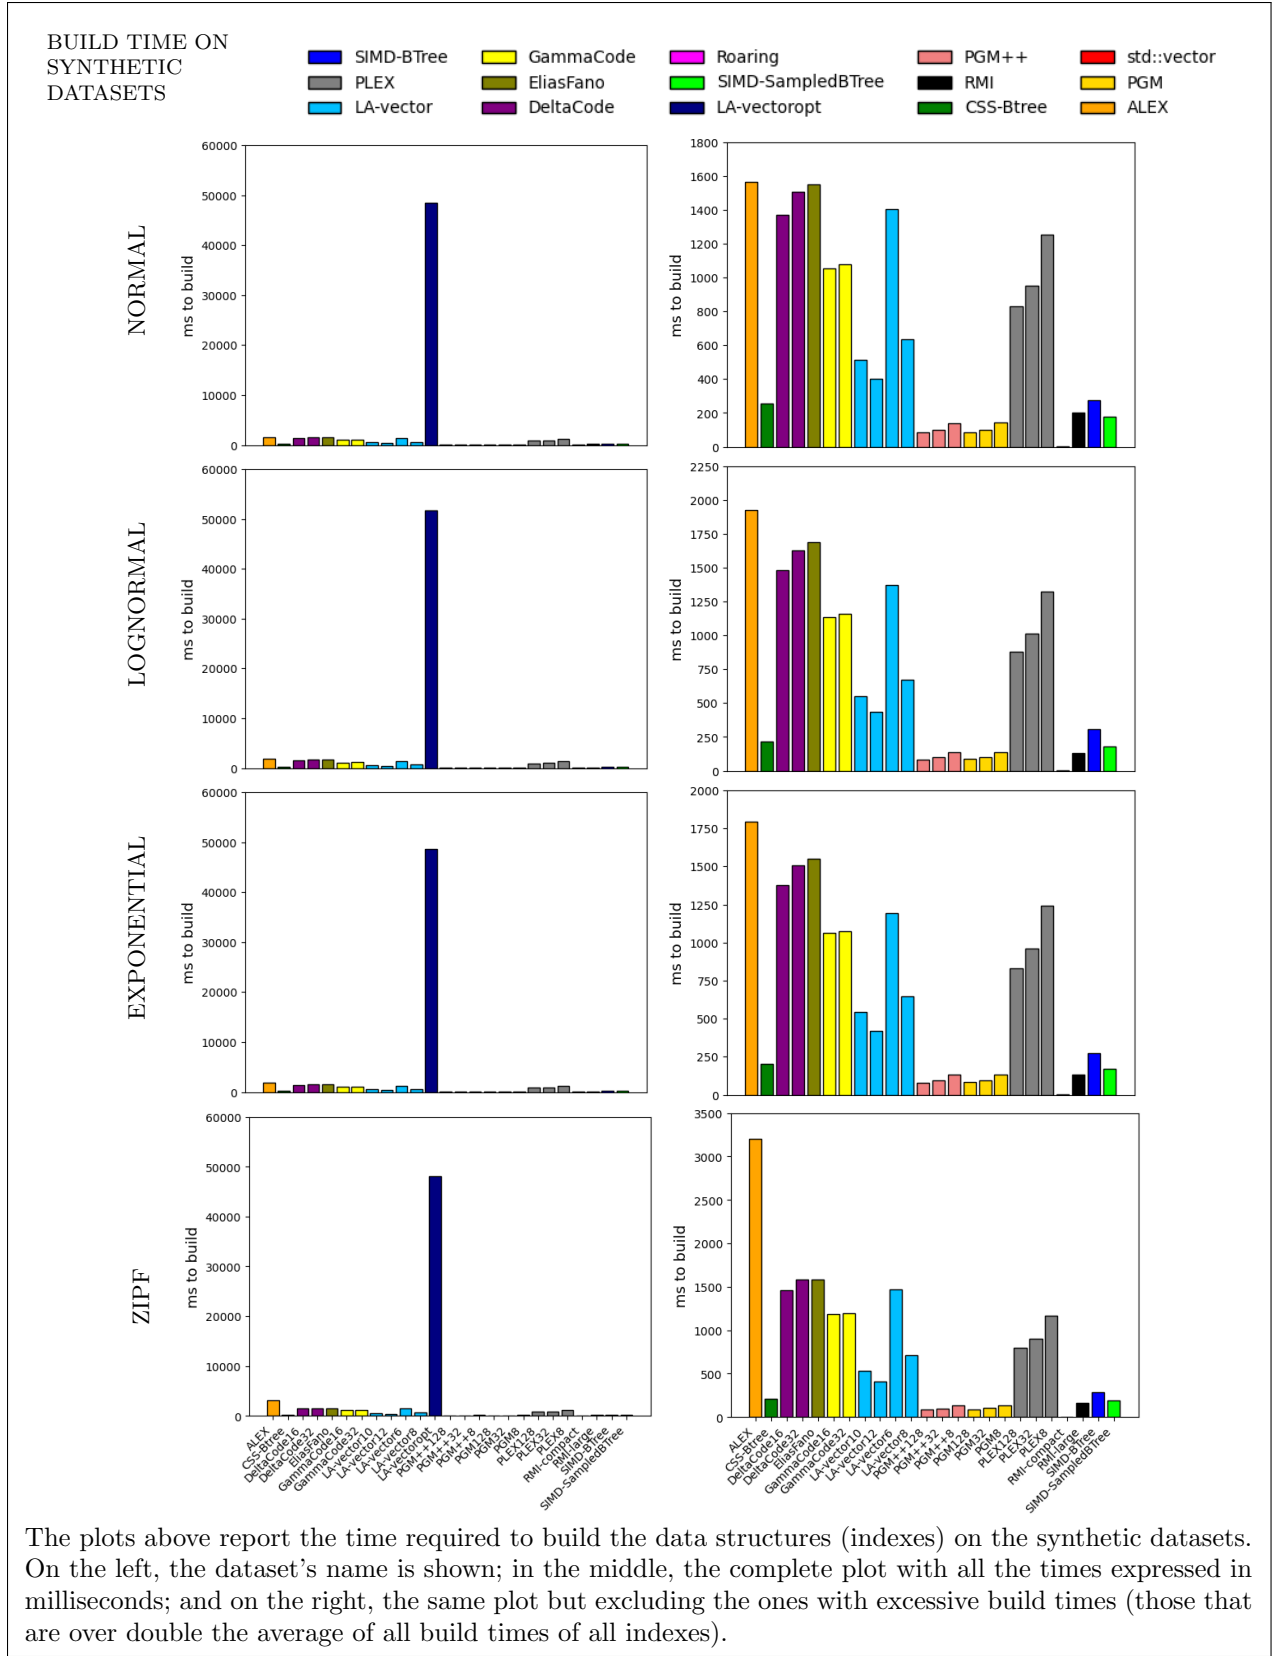


Figure 14: Build time on synthetic datasets

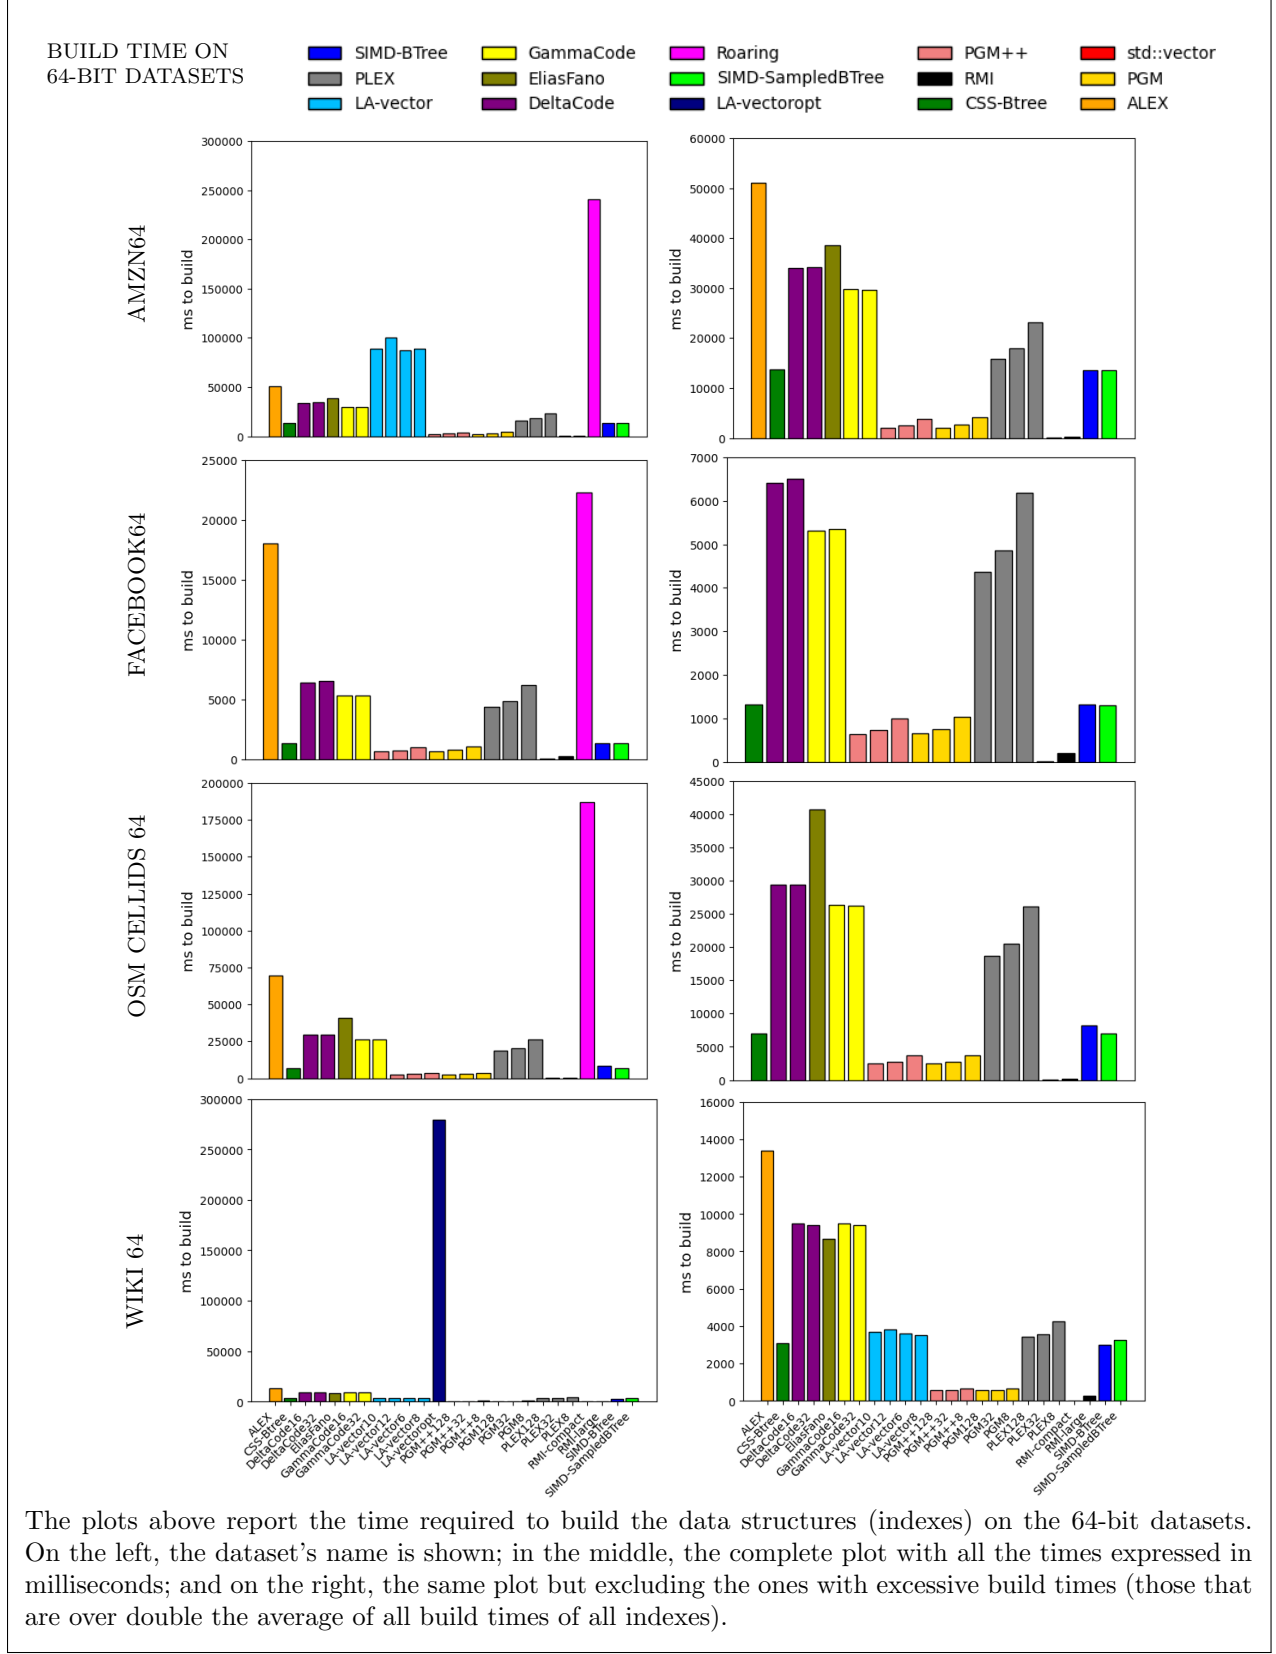
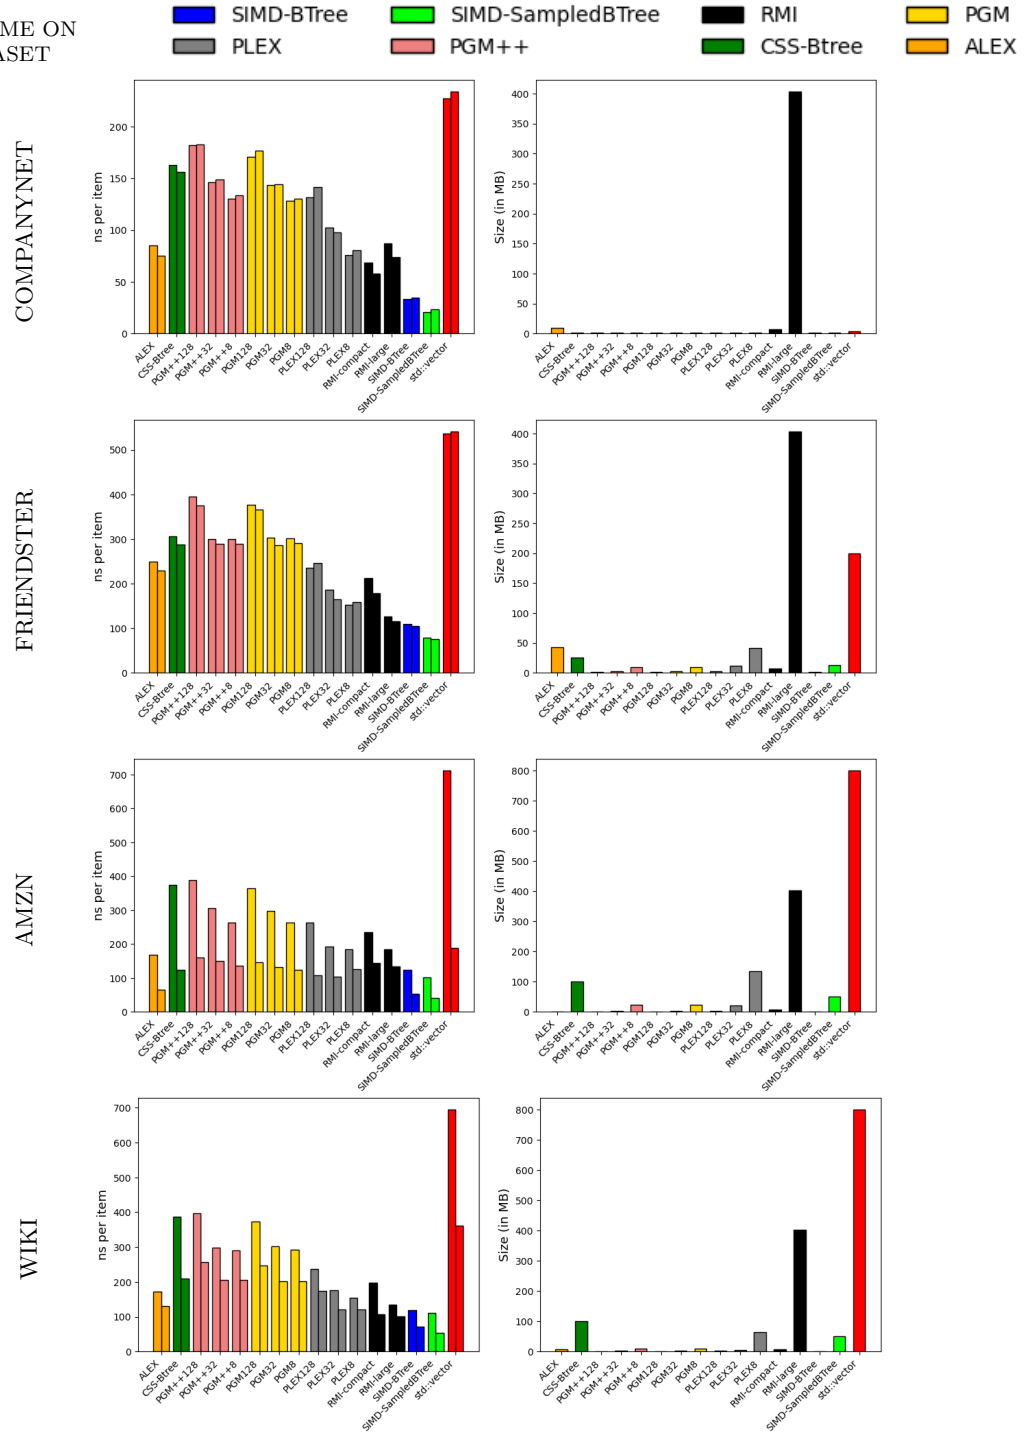


Figure 15: Build time on 64-bit datasets

TRADITIONAL/LEARNED
INDEXES::
SEARCH TIME ON
REAL DATASET





The performances of traditional and learned indexes on real-world datasets are represented in two plots. The leftmost plot shows the time of search (in ns) on each index; while the right plot shows the space occupied by each index in MBs (without considering the space occupied by the data to be stored). The left plot shows two bars for each data structure. The left/right one shows the average time to search for an existing/missing item in the collection.

Figure 16: Average time for pointwise queries on traditional and learned indexes, built on real-world datasets.



TRADITIONAL/LEARNED
INDEXES::
SEARCH TIME ON
SYNTHETIC DATASET

■ SIMD-BTree ■ SIMD-SampledBTree ■ RMI ■ PGM
■ PLEX ■ PGM++ ■ CSS-Btree ■ ALEX



NORMAL





LOGNORMAL



EXPONENTIAL



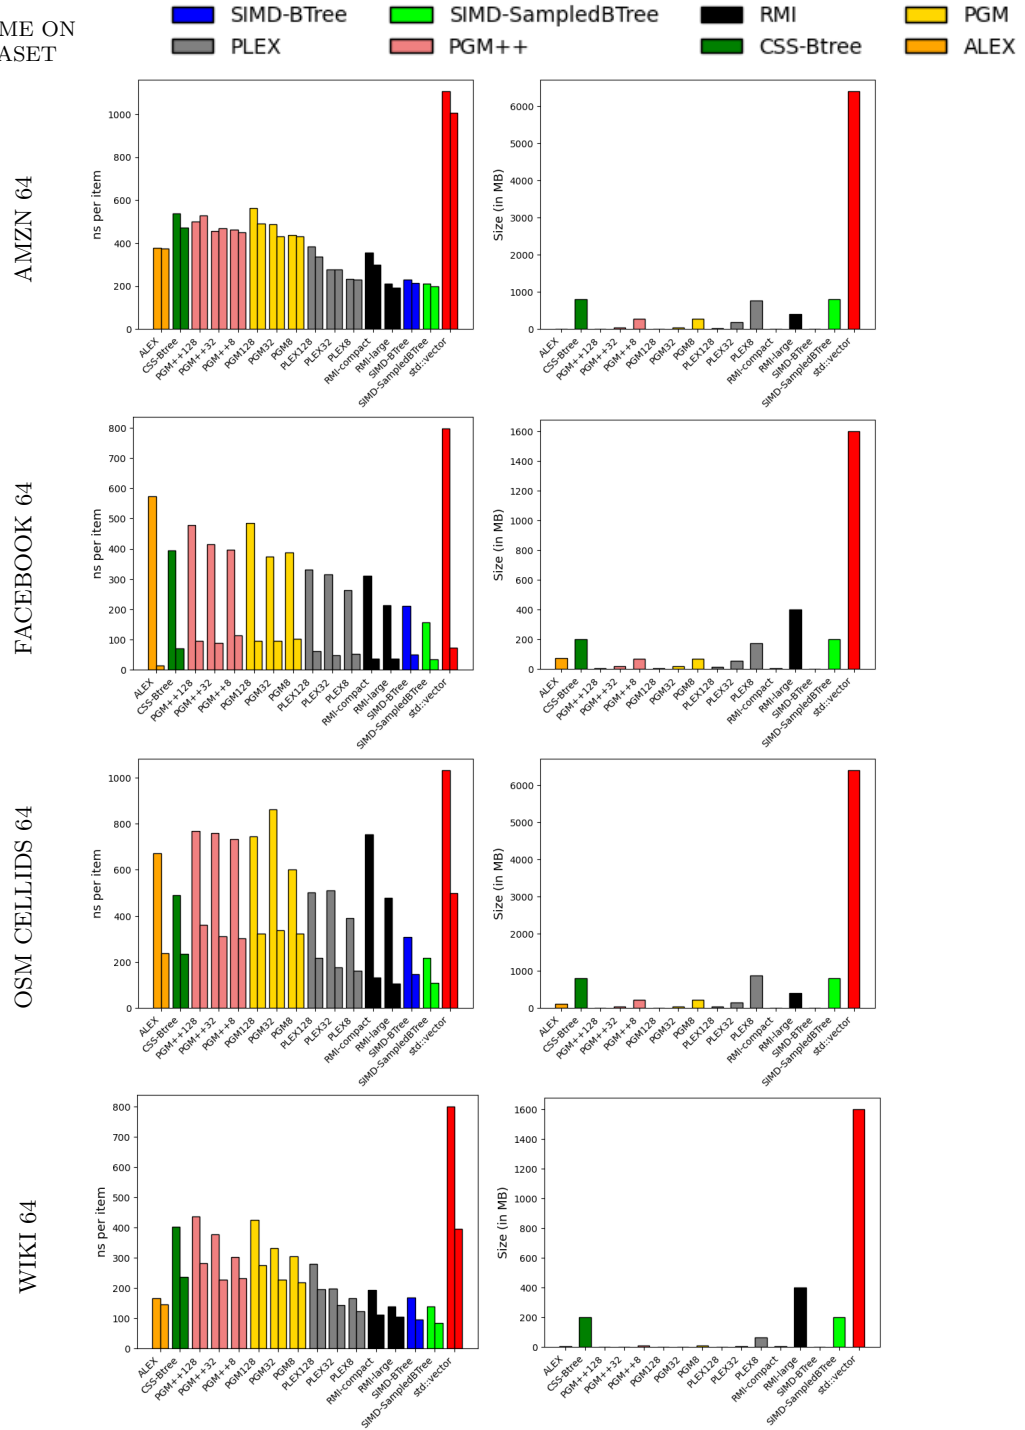
ZIPF



The performance of traditional and learned indexes on synthetic datasets are represented in two plots. The leftmost plot shows the time of search (in ns) on each index; while the right plot shows the space occupied by each index in MBs (without considering the space occupied by the data to be stored). The left plot shows two bars for each data structure. The left/right one shows the average time to search for an existing/missing item in the collection.

Figure 17: Average time for pointwise queries on traditional and learned indexes, built on synthetic datasets.

TRADITIONAL/LEARNED
INDEXES::
SEARCH TIME ON
64-BIT DATASET




The performances of traditional and learned indexes on 64-bit datasets are represented in two plots. The leftmost plot shows the time of search (in ns) on each index; while the right plot shows the space occupied by each index in MBs (without considering the space occupied by the data to be stored). The left plot shows two bars for each data structure. The left/right one shows the average time to search for an existing/missing item in the collection.

Figure 18: Average time for pointwise queries on traditional and learned indexes, built on 64-bit datasets.

COMPRESSED INDEXES::
SEARCH TIME ON
REAL DATASET

LA-vector EliasFano LA-vectoropt DeltaCode
GammaCode Roaring std::vector

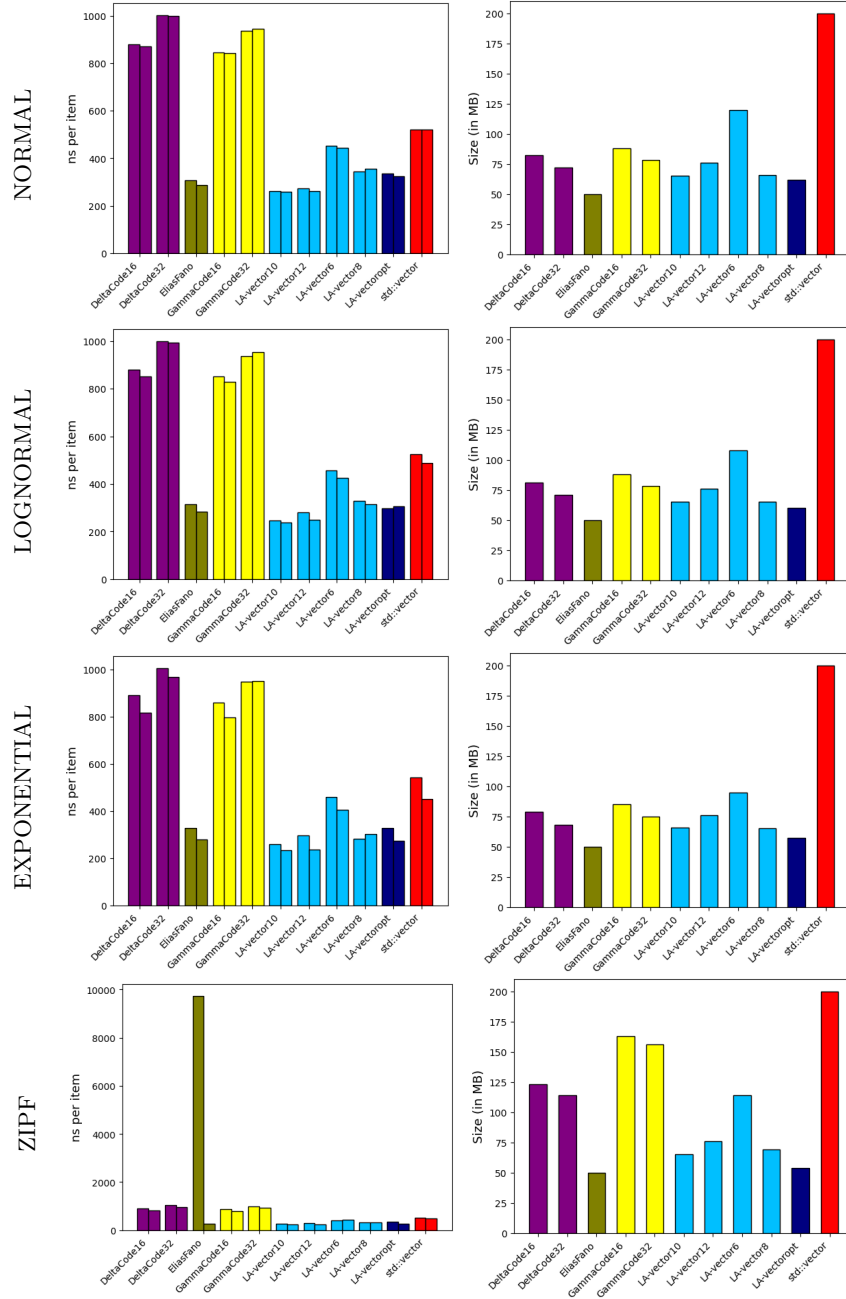


The performances of compressed indexes on real-world datasets are represented in two plots. The leftmost plot shows the time of search (in ns) on each index; while the right plot shows the space occupied by each index in MBs (without considering the space occupied by the data to be stored). The left plot shows two bars for each data structure. The left/right one shows the average time to search for an existing/missing item in the collection.

Figure 19: Average time for pointwise queries on compressed indexes, built on real-world datasets.

COMPRESSED INDEXES::
SEARCH TIME ON
SYNTHETIC DATASET


LA-vector EliasFano LA-vectoropt DeltaCode
GammaCode Roaring std::vector



The performances of compressed indexes on synthetic datasets are represented in two plots. The leftmost plot shows the time of search (in ns) on each index; while the right plot shows the space occupied by each index in MBs (without considering the space occupied by the data to be stored). The left plot shows two bars for each data structure. The left/right one shows the average time to search for an existing/missing item in the collection.

Figure 20: Average time for pointwise queries on compressed indexes, built on synthetic datasets.

COMPRESSED INDEXES::
SEARCH TIME ON
64-BIT DATASET




The performances of compressed indexes on 64-bit datasets are represented in two plots. The leftmost plot shows the time of search (in ns) on each index; while the right plot shows the space occupied by each index in MBs (without considering the space occupied by the data to be stored). The left plot shows two bars for each data structure. The left/right one shows the average time to search for an existing/missing item in the collection.

Figure 21: Average time for pointwise queries on compressed indexes, built on 64-bit datasets.

COMPRESSED INDEXES::
SCAN TIME

LA-vector EliasFano LA-vectoropt DeltaCode
GammaCode Roaring std::vector



The plots above show the performance of compressed indexes in time and space relative to *range queries*, where starting points are randomly sampled, and the width of the scan is set to 10, 100, 1K, and 10K. The plot on the left shows the time (in ns) required for every interrogation, describing the average time required per access with $x = 10, 100, 1K, 10K$. The plot on the right shows the space occupied by each compressed index (in MB).

Figure 22: Average time (in ns) for range queries on compressed indexes, on 4 real-world datasets.

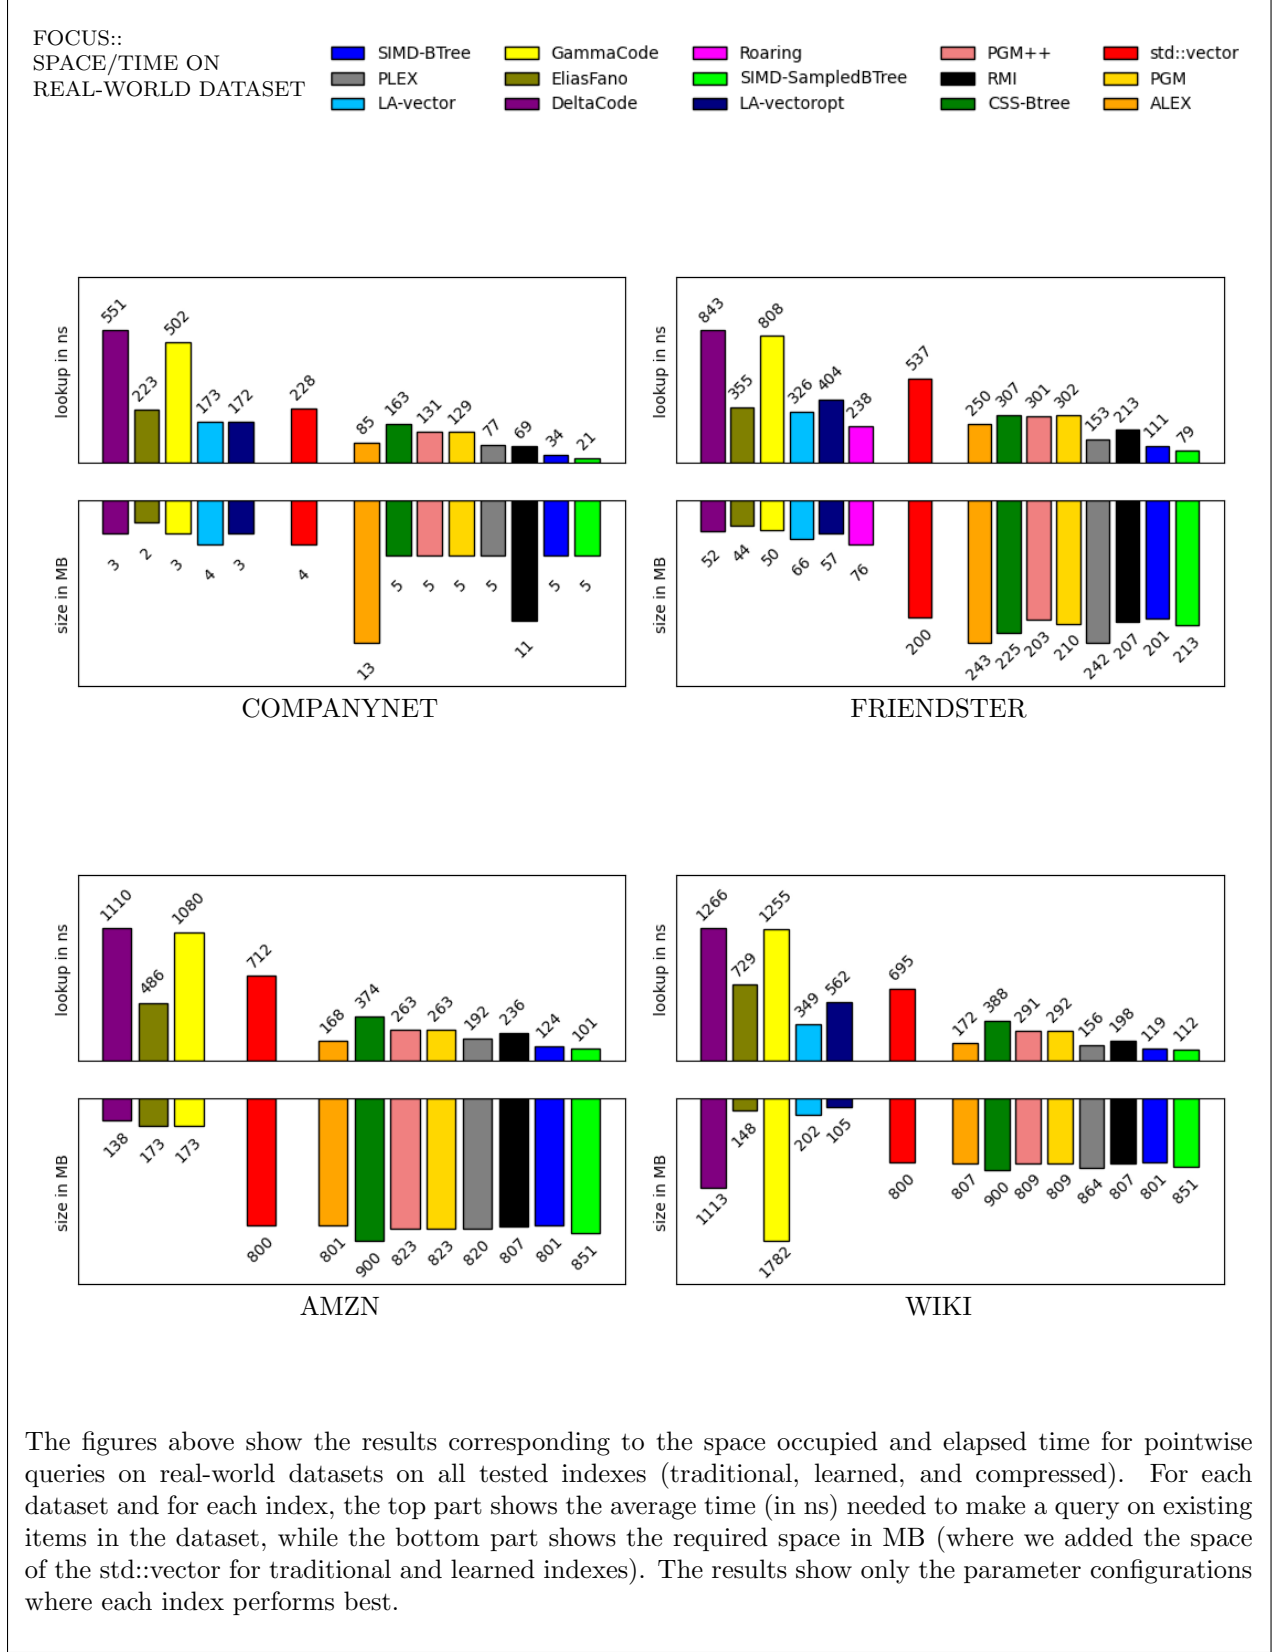


Figure 23: Recap space/time plots for pointwise queries on real-world datasets.

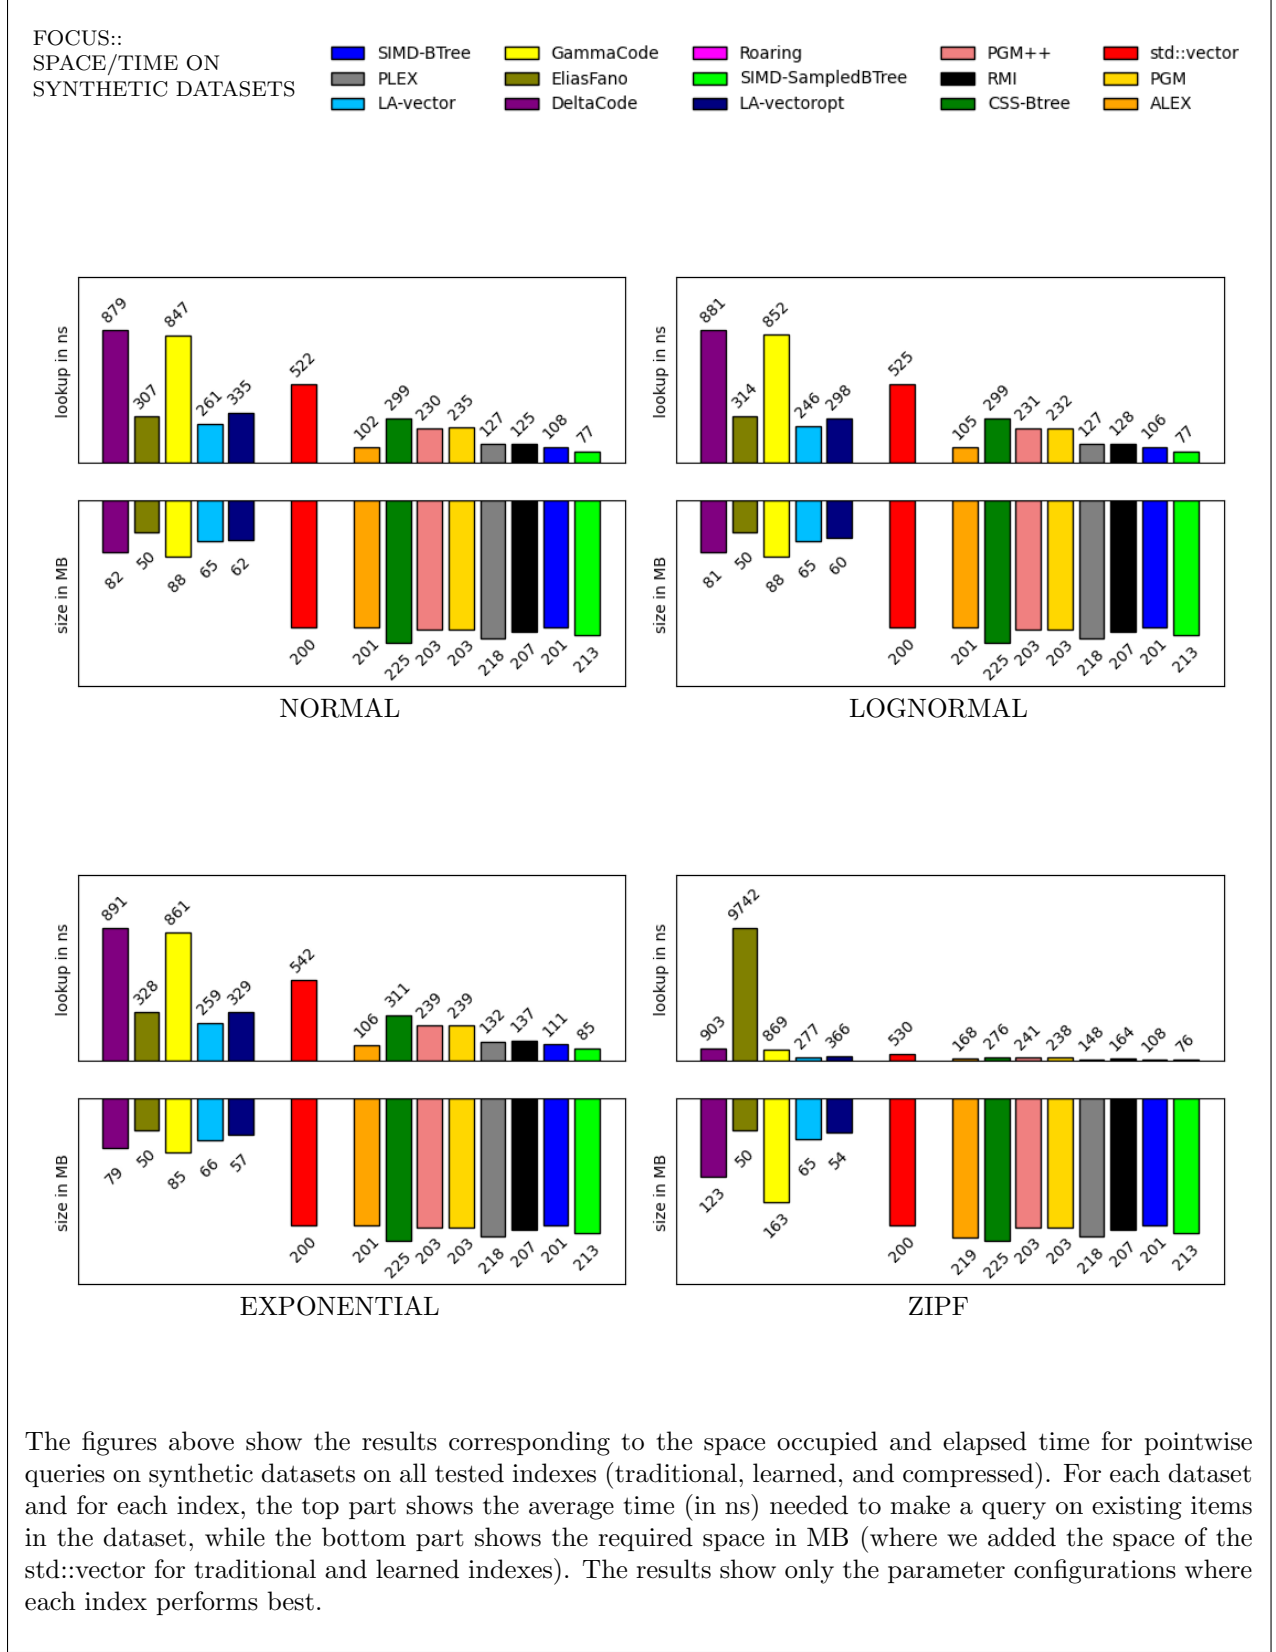


Figure 24: Recap space/time plots for pointwise queries on synthetic datasets.

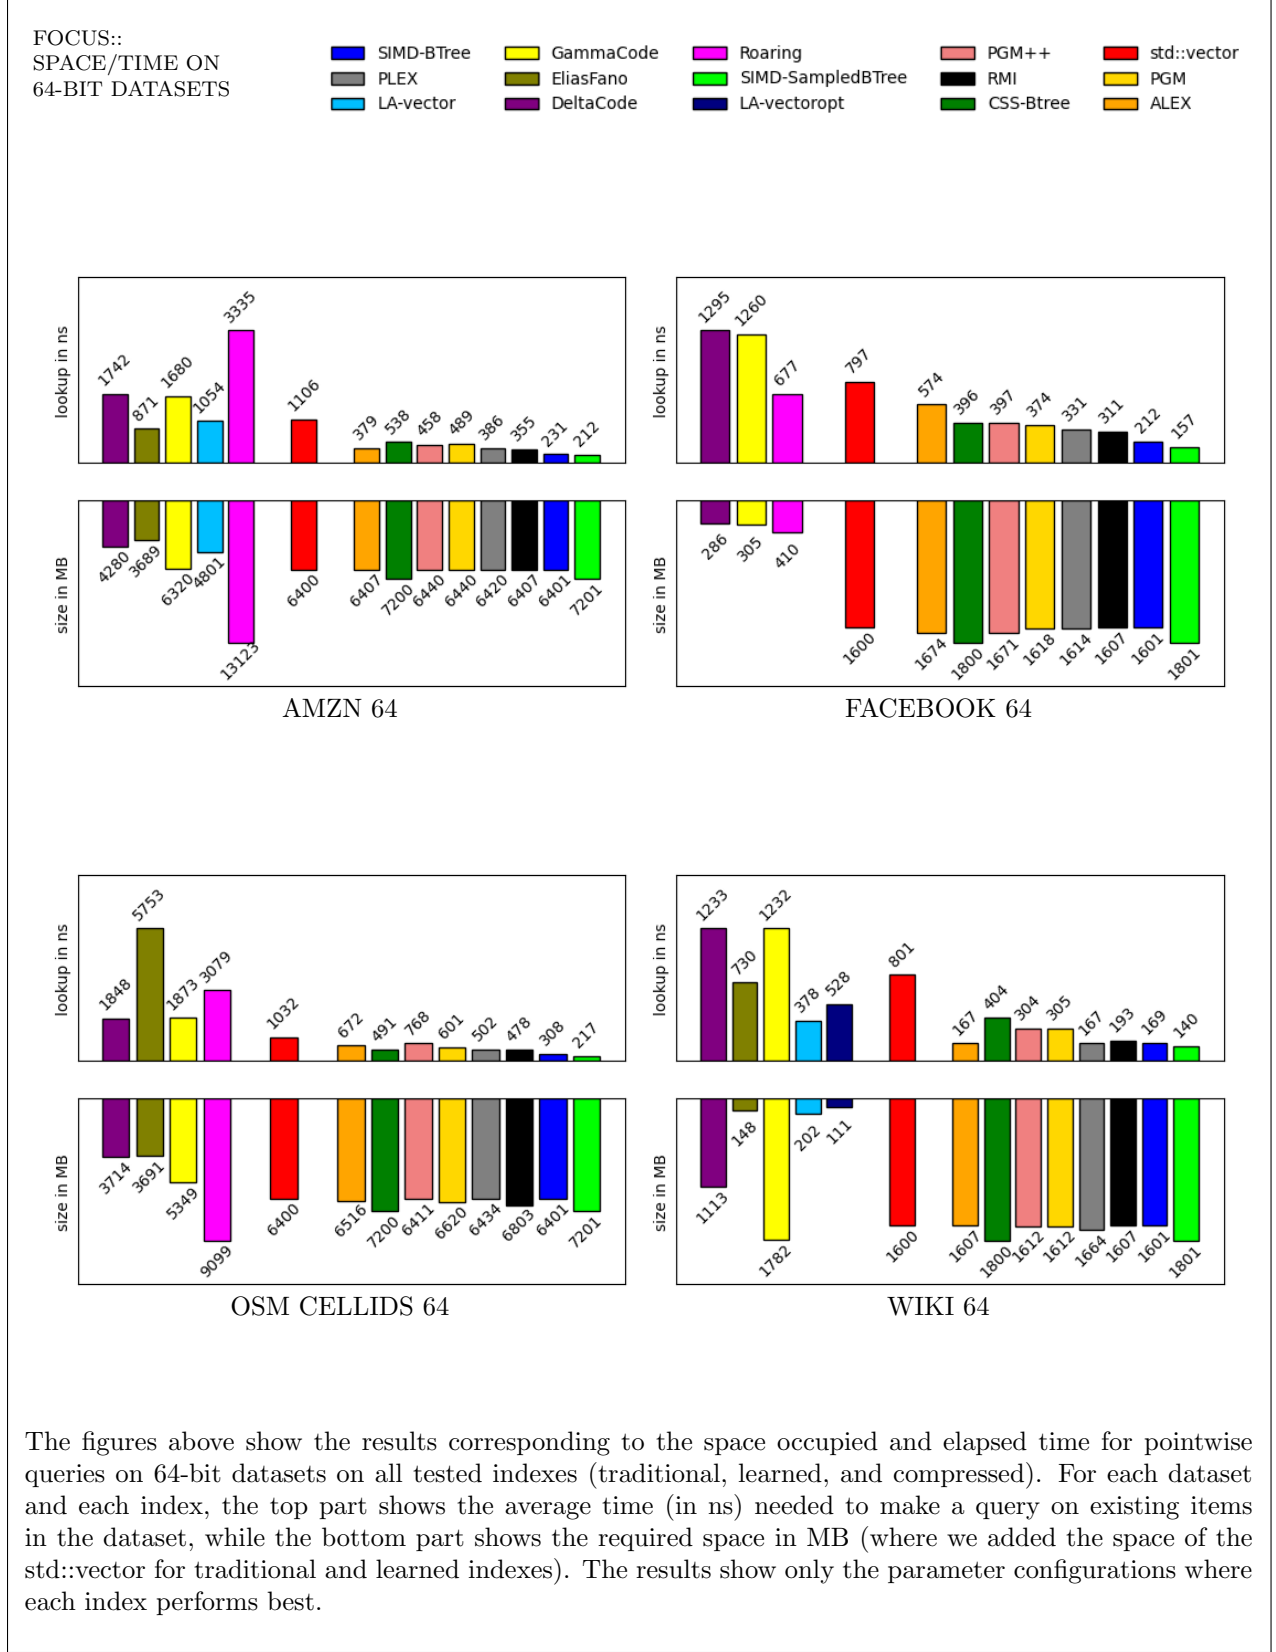


Figure 25: Recap space/time plots for pointwise queries on 64-bit datasets.

FOCUS:: PARETO CURVE
ON TRADITIONAL
AND LEARNED INDEXES
ON REAL-WORLD
DATASETS



ALEX



PGM++8



PLEX128



RMI-compact



CSS-Btree



PGM128



PLEX32



SIMD-BTree



PGM++128



PGM32



PLEX8



PGM8



PGM++32




PGM8




PLEX8




SIMD-SampledBTree




COMPANYNET



FRIENDSTER



AMZN



WIKI

The plots recap the experimental results about occupied space/time needed for pointwise queries for traditional and learned indexes on real-world datasets. For these datasets, for each index, and each tested configuration we plot the extra space occupied (in MB) on the x-axis, and the time (in ns) per pointwise query on existing items in the datasets.

Additionally, a black line shows the Pareto frontier for traditional/learned indexes. The indexes that sit on top of the Pareto frontier offer the best space-time trade-off. We avoided plotting “RMI-large” because of its excessive space occupancy (roughly 400 MB on each dataset). “RMI-large” was not one of the Pareto-optimal configurations.

Figure 26: Pareto Frontier: Traditional and Learned Indexes on real-world datasets

FOCUS:: PARETO CURVE
ON TRADITIONAL
AND LEARNED INDEXES
ON SYNTHETIC
DATASETS



ALEX



PGM++8



PLEX128



RMI-compact



CSS-Btree



PGM128



PLEX32



SIMD-BTree



PGM++128



PGM32



PLEX8







SIMD-SampledBTree



PGM++32



PGM8



The plots recap the experimental results about occupied space/time needed for pointwise queries for traditional and learned indexes on synthetic datasets. For these datasets, for each index, and each tested configuration we plot the extra space occupied (in MB) on the x-axis, and the time (in ns) per pointwise query on existing items in the datasets.

Additionally, a black line shows the Pareto frontier for traditional/learned indexes. The indexes that sit on top of the Pareto frontier offer the best space-time trade-off. We avoided plotting “RMI-large” because of its excessive space occupancy (roughly 400 MB on each dataset). “RMI-large” was not one of the Pareto-optimal configurations.

Figure 27: Pareto Frontier: Traditional and Learned Indexes on synthetic datasets

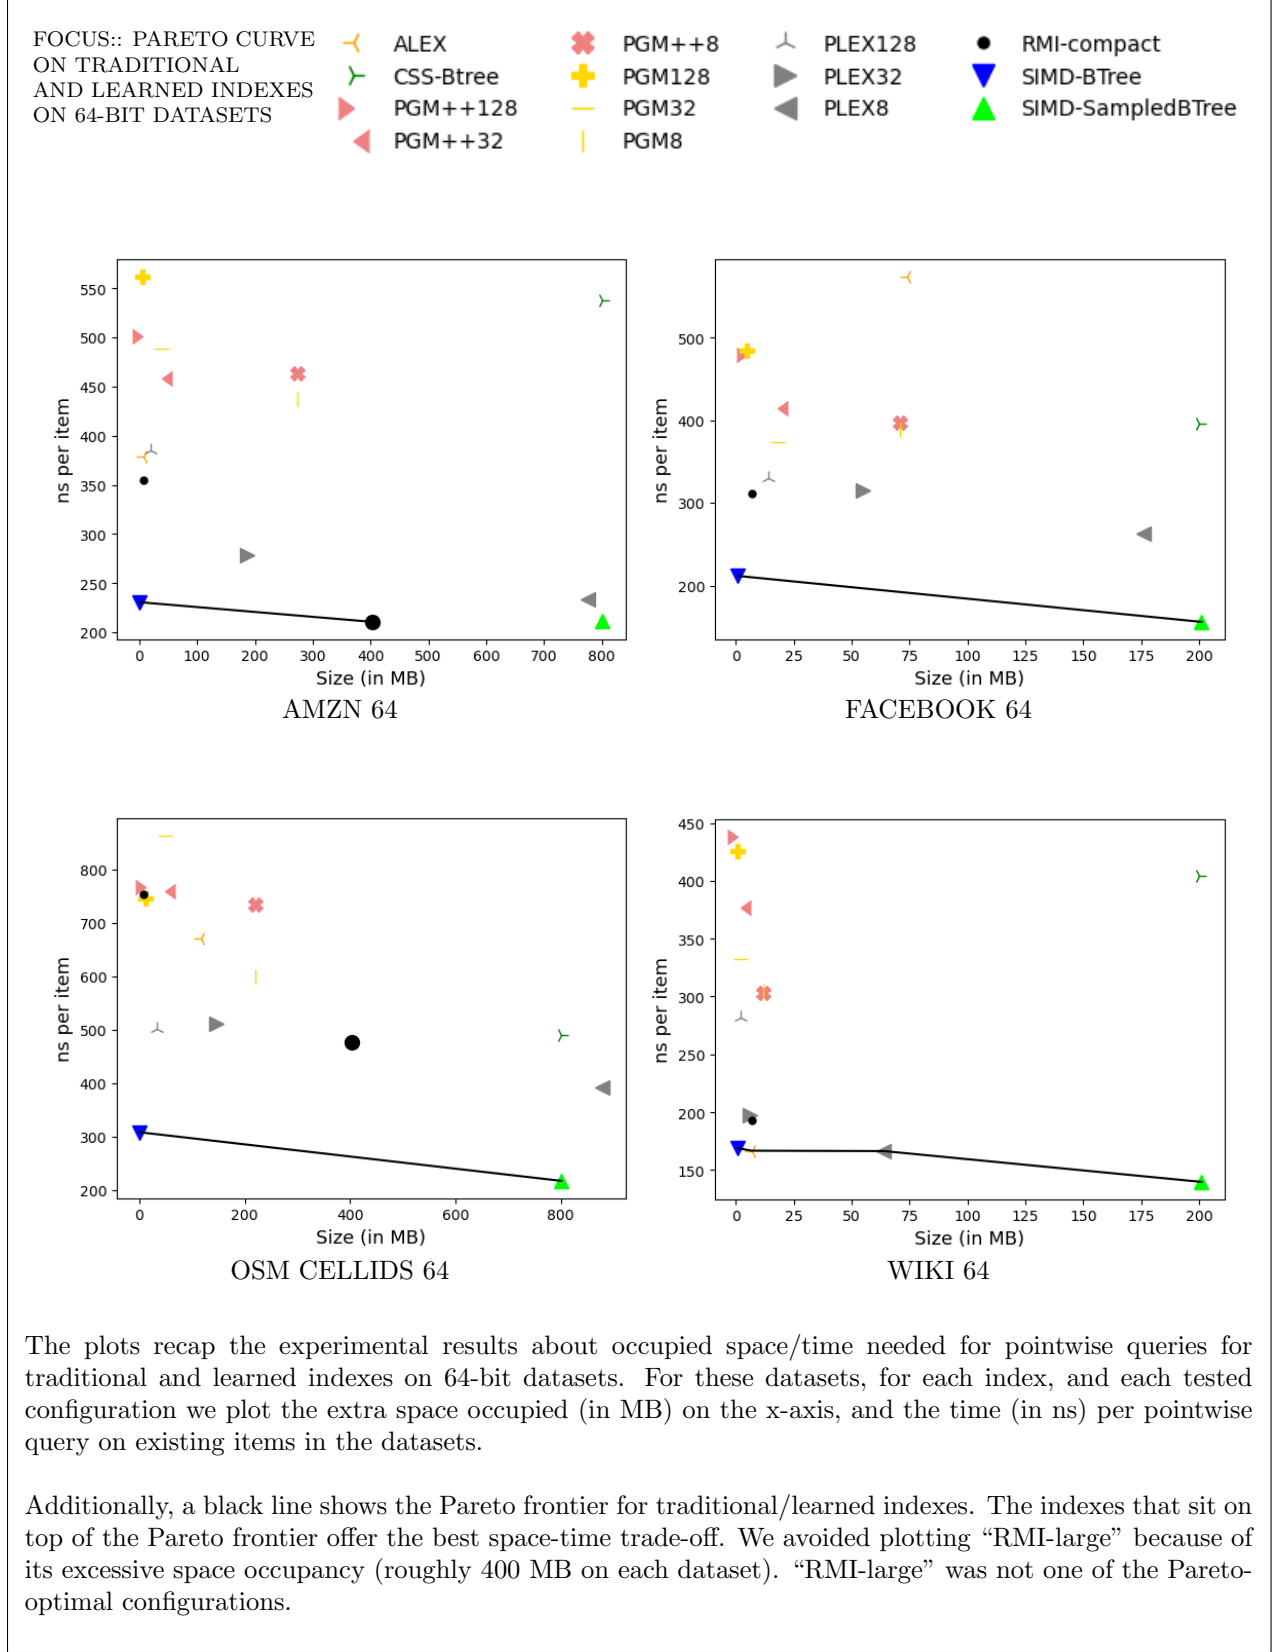


Figure 28: Pareto Frontier: Traditional and Learned Indexes on 64-bit datasets

FOCUS:: PARETO

CURVE ON
COMPRESSED INDEXES
ON REAL-WORLD
DATASETS



DeltaCode16



GammaCode16



LA-vector12



LA-vectoropt



DeltaCode32



GammaCode32



LA-vector6



std::vector




EliasFano




LA-vector10




LA-vector8




COMPANYNET



FRIENDSTER



AMZN



WIKI

The plots recap the experimental results about occupied space/time needed for pointwise queries for compressed indexes on real-world datasets. For these datasets, for each compressed index, and each tested configuration we plot the extra space occupied (in MB) on the x-axis, and the time (in ns) per pointwise query on existing items in the datasets.

Additionally, a black line shows the Pareto frontier for traditional/learned indexes. The indexes that sit on top of the Pareto frontier offer the best space-time trade-off.

Figure 29: Pareto Frontier: Compressed Indexes on real-world datasets

FOCUS:: PARETO


CURVE ON
COMPRESSED INDEXES
ON SYNTHETIC
DATASETS

DeltaCode16
DeltaCode32
EliasFano


GammaCode16
GammaCode32
LA-vector10

LA-vector12
LA-vector6
LA-vector8


LA-vectoropt
std::vector




NORMAL



LOGNORMAL



EXPONENTIAL



ZIPF

The plots recap the experimental results about occupied space/time needed for pointwise queries for compressed indexes on synthetic datasets. For these datasets, for each compressed index, and each tested configuration we plot the extra space occupied (in MB) on the x-axis, and the time (in ns) per pointwise query on existing items in the datasets.

Additionally, a black line shows the Pareto frontier for traditional/learned indexes. The indexes that sit on top of the Pareto frontier offer the best space-time trade-off.

Figure 30: Pareto Frontier: Compressed Indexes on synthetic datasets

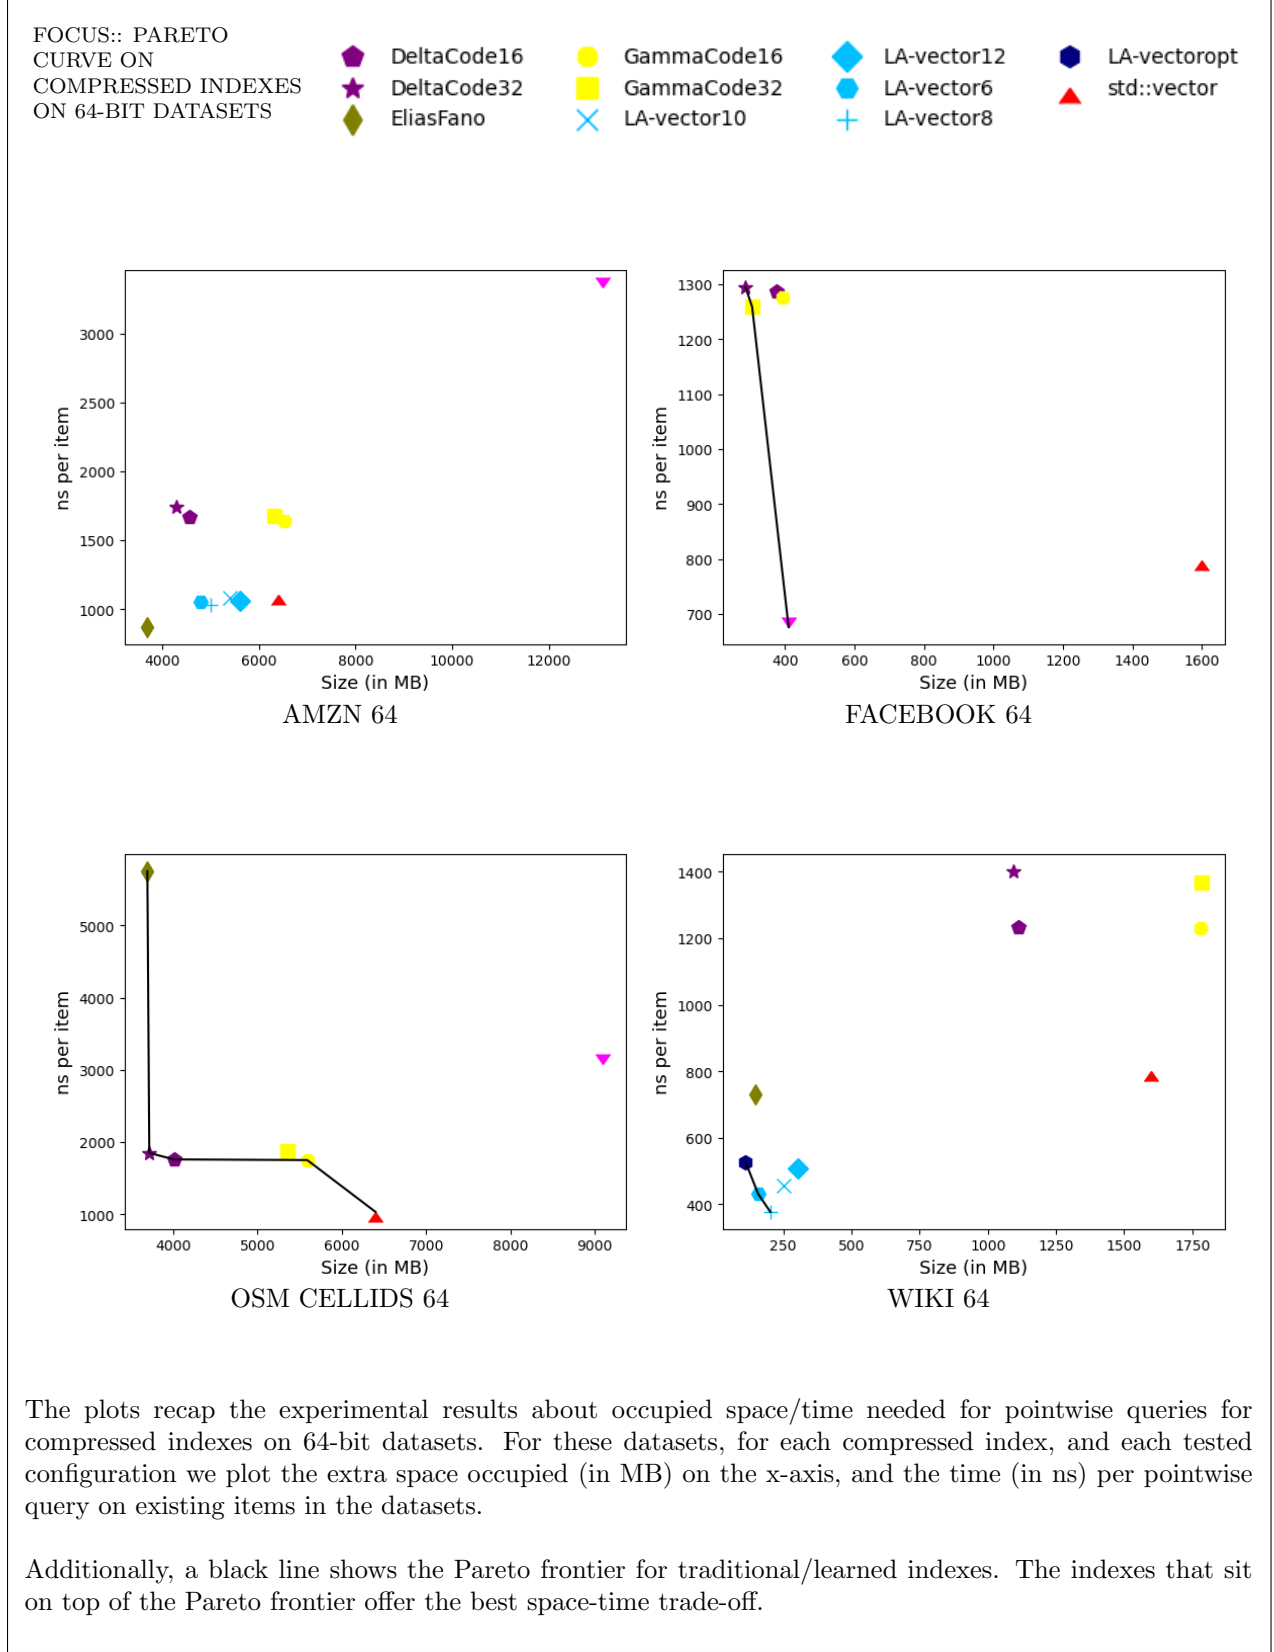


Figure 31: Pareto Frontier: Compressed Indexes on 64-bit datasets

| index | lookup existing (ns) | lookup missing (ns) | space (MB) | build (ms) |
|-------------------|----------------------|---------------------|------------|-------------|
| ALEX | 84.795 | 75.2779 | 9 | 73.757 |
| CSS-Btree | 162.639 | 156.072 | 1 | 3.47303 |
| DeltaCode16 | 550.163 | 550.164 | 3 | 28.4947 |
| DeltaCode32 | 801.402 | 801.772 | 3 | 31.9948 |
| EliasFano | 222.173 | 190.123 | 2 | 34.5756 |
| GammaCode16 | 501.978 | 497.942 | 3 | 21.8515 |
| GammaCode32 | 718.739 | 716.487 | 3 | 22.4363 |
| LA-vector10 | 174.389 | 163.978 | 5 | 49.1686 |
| LA-vector12 | 172.844 | 158.693 | 4 | 31.8208 |
| LA-vector6 | 177.356 | 179.273 | 7 | 55.6362 |
| LA-vector8 | 177.742 | 165.176 | 6 | 50.2594 |
| LA-vectoropt | 171.206 | 174.05 | 3 | 959.086 |
| PGM++128 | 181.748 | 182.408 | 1 | 2.34623 |
| PGM++32 | 145.778 | 148.903 | 1 | 2.33863 |
| PGM++8 | 130.347 | 133.17 | 1 | 3.36725 |
| PGM128 | 170.391 | 176.688 | 1 | 2.39936 |
| PGM32 | 143.277 | 144.125 | 1 | 2.3582 |
| PGM8 | 128.272 | 129.927 | 1 | 3.40615 |
| PLEX128 | 131.528 | 141.461 | 1 | 18.8695 |
| PLEX32 | 102.084 | 97.9293 | 1 | 20.4044 |
| PLEX8 | 76.0505 | 80.3045 | 1 | 24.9302 |
| RMI-compact | 68.2063 | 57.9312 | 7 | 3.14833 |
| RMI-large | 87.104 | 73.6292 | 403 | 138.054 |
| SIMD-BTree | 33.1707 | 34.8978 | 1 | 4.20203 |
| SIMD-SampledBTree | 20.9854 | 23.1505 | 1 | 2.84402 |
| std::vector | 227.127 | 233.441 | 4 | 0.000196509 |

Table 1: Tabular data: companynet dataset

| index | lookup existing (ns) | lookup missing (ns) | space (MB) | build (ms) |
|-------------------|----------------------|---------------------|------------|-------------|
| ALEX | 171.216 | 130.978 | 7 | 12250.3 |
| CSS-Btree | 387.456 | 209.897 | 100 | 1706.93 |
| DeltaCode16 | 1265.27 | 766.646 | 1113 | 8585.24 |
| DeltaCode32 | 1406.5 | 951.687 | 1095 | 8918.44 |
| EliasFano | 728.747 | 251.374 | 148 | 7772.2 |
| GammaCode16 | 1254.4 | 743.391 | 1782 | 8558.24 |
| GammaCode32 | 1395.82 | 923.777 | 1785 | 8523.53 |
| LA-vector10 | 419.032 | 237.91 | 251 | 1924.89 |
| LA-vector12 | 462.707 | 258.213 | 301 | 1744.62 |
| LA-vector6 | 380.864 | 315.315 | 154 | 2069.24 |
| LA-vector8 | 348.619 | 241.481 | 202 | 1795.67 |
| LA-vectoropt | 561.527 | 285.196 | 105 | 197674 |
| PGM++128 | 396.57 | 256.087 | 1 | 486.691 |
| PGM++32 | 298.428 | 204.902 | 2 | 482.605 |
| PGM++8 | 290.848 | 205.049 | 9 | 564.017 |
| PGM128 | 374.152 | 246.121 | 1 | 492.813 |
| PGM32 | 302.538 | 201.409 | 2 | 493.626 |
| PGM8 | 291.528 | 201.788 | 9 | 569.178 |
| PLEX128 | 237.739 | 174.58 | 2 | 3024.15 |
| PLEX32 | 176.581 | 121.453 | 6 | 3128.08 |
| PLEX8 | 155.252 | 120.38 | 64 | 3749.19 |
| RMI-compact | 197.538 | 106.451 | 7 | 2.84652 |
| RMI-large | 135.004 | 101.825 | 403 | 235.648 |
| SIMD-BTree | 118.831 | 71.6056 | 1 | 1597.51 |
| SIMD-SampledBTree | 111.367 | 53.0724 | 51 | 1569.05 |
| std::vector | 694.065 | 360.534 | 800 | 0.000316836 |

Table 2: Tabular data: wiki uint32 dataset

| index | lookup existing (ns) | lookup missing (ns) | space (MB) | build (ms) |
|-------------------|----------------------|---------------------|------------|-------------|
| ALEX | 249.282 | 228.425 | 43 | 3246.62 |
| CSS-Btree | 306.611 | 287.862 | 25 | 202.872 |
| DeltaCode16 | 842.925 | 838.148 | 52 | 1304.31 |
| DeltaCode32 | 907.486 | 935.238 | 41 | 1437.65 |
| EliasFano | 354.424 | 268.447 | 44 | 1888.67 |
| GammaCode16 | 807.403 | 801.357 | 50 | 1035.2 |
| GammaCode32 | 873.273 | 882.188 | 39 | 1042.19 |
| LA-vector10 | 325.142 | 276.805 | 66 | 526.315 |
| LA-vector12 | 345.036 | 313.693 | 76 | 434.242 |
| LA-vector6 | 402.512 | 323 | 62 | 811.986 |
| LA-vector8 | 350.417 | 310.847 | 60 | 576.623 |
| LA-vectoropt | 403.264 | 384.947 | 57 | 52328.7 |
| PGM++128 | 394.567 | 375.317 | 1 | 131.291 |
| PGM++32 | 300.058 | 288.891 | 3 | 157.844 |
| PGM++8 | 300.258 | 288.922 | 10 | 204.42 |
| PGM128 | 376.135 | 365.009 | 1 | 132.605 |
| PGM32 | 302.368 | 286.235 | 3 | 159.166 |
| PGM8 | 301.493 | 290.429 | 10 | 207.479 |
| PLEX128 | 234.662 | 246.05 | 3 | 935.696 |
| PLEX32 | 185.963 | 165.233 | 12 | 1099.88 |
| PLEX8 | 152.585 | 158.389 | 42 | 1453.34 |
| RMI-compact | 212.397 | 177.803 | 7 | 1.70732 |
| RMI-large | 126.462 | 116.152 | 403 | 210.646 |
| Roaring | 237.821 | 320.745 | 76 | 2254.53 |
| SIMD-BTree | 110.103 | 104.483 | 1 | 292.077 |
| SIMD-SampledBTree | 78.882 | 74.924 | 13 | 169.295 |
| std::vector | 536.189 | 540.293 | 200 | 0.000377931 |

Table 3: Tabular data: friendster dataset

| index | lookup existing (ns) | lookup missing (ns) | space (MB) | build (ms) |
|-------------------|----------------------|---------------------|------------|-------------|
| ALEX | 105.054 | 85.9853 | 1 | 1793.16 |
| CSS-Btree | 310.346 | 252.282 | 25 | 205.264 |
| DeltaCode16 | 890.872 | 816.098 | 79 | 1375.91 |
| DeltaCode32 | 1006.65 | 969.474 | 68 | 1508.29 |
| EliasFano | 327.316 | 280.85 | 50 | 1548.45 |
| GammaCode16 | 860.252 | 798.242 | 85 | 1065.69 |
| GammaCode32 | 949.054 | 950.769 | 75 | 1072.15 |
| LA-vector10 | 258.936 | 234.123 | 66 | 541.29 |
| LA-vector12 | 297.423 | 236.388 | 76 | 419.269 |
| LA-vector6 | 459.798 | 405.767 | 95 | 1195.47 |
| LA-vector8 | 281.28 | 303.526 | 65 | 648.813 |
| LA-vectoropt | 328.516 | 273.581 | 57 | 48586.3 |
| PGM++128 | 333.256 | 273.656 | 1 | 80.6546 |
| PGM++32 | 268.757 | 236.751 | 1 | 93.8582 |
| PGM++8 | 238.733 | 211.15 | 3 | 130.359 |
| PGM128 | 296.382 | 260.838 | 1 | 81.6738 |
| PGM32 | 261.463 | 224.831 | 1 | 95.6289 |
| PGM8 | 238.67 | 207.413 | 3 | 132.733 |
| PLEX128 | 218.257 | 196.585 | 1 | 831.188 |
| PLEX32 | 164.447 | 136.248 | 2 | 958.108 |
| PLEX8 | 131.953 | 121.509 | 18 | 1244.44 |
| RMI-compact | 136.908 | 110.583 | 7 | 2.86306 |
| RMI-large | 100.939 | 101.656 | 403 | 130.797 |
| SIMD-BTree | 110.86 | 85.3157 | 1 | 273.07 |
| SIMD-SampledBTree | 84.1393 | 61.7075 | 13 | 169.157 |
| std::vector | 541.996 | 450.289 | 200 | 0.000363961 |

Table 4: Tabular data: exponential dataset

| index | lookup existing (ns) | lookup missing (ns) | space (MB) | build (ms) |
|-------------------|----------------------|---------------------|------------|-------------|
| ALEX | 104.909 | 91.3408 | 1 | 1925.96 |
| CSS-Btree | 298.479 | 280.769 | 25 | 218.561 |
| DeltaCode16 | 880.698 | 851.485 | 81 | 1481.7 |
| DeltaCode32 | 999.842 | 994.071 | 71 | 1628.53 |
| EliasFano | 313.658 | 283.585 | 50 | 1691 |
| GammaCode16 | 851.016 | 828.575 | 88 | 1133.47 |
| GammaCode32 | 937.174 | 953.315 | 78 | 1160.45 |
| LA-vector10 | 245.619 | 237.983 | 65 | 549.706 |
| LA-vector12 | 279.572 | 250.127 | 76 | 432.959 |
| LA-vector6 | 456.651 | 425.821 | 108 | 1374.75 |
| LA-vector8 | 329.223 | 314.825 | 65 | 673.096 |
| LA-vectoropt | 297.824 | 307.213 | 60 | 51735.8 |
| PGM++128 | 319.303 | 272.324 | 1 | 84.8645 |
| PGM++32 | 256.398 | 223.029 | 1 | 100.779 |
| PGM++8 | 230.012 | 204.547 | 3 | 139.745 |
| PGM128 | 287.646 | 253.085 | 1 | 86.8909 |
| PGM32 | 246.915 | 219.002 | 1 | 102.837 |
| PGM8 | 231.244 | 202.007 | 3 | 139.112 |
| PLEX128 | 209.419 | 209.191 | 1 | 881.072 |
| PLEX32 | 158.246 | 145.083 | 2 | 1011.73 |
| PLEX8 | 126.922 | 128.496 | 18 | 1326.59 |
| RMI-compact | 127.217 | 118.769 | 7 | 3.23318 |
| RMI-large | 116.657 | 121.766 | 403 | 132.918 |
| SIMD-BTree | 105.188 | 104.163 | 1 | 307.019 |
| SIMD-SampledBTree | 76.7902 | 71.7349 | 13 | 182.324 |
| std::vector | 524.067 | 488.268 | 200 | 0.000275858 |

Table 5: Tabular data: lognormal dataset

| index | lookup existing (ns) | lookup missing (ns) | space (MB) | build (ms) |
|-------------------|----------------------|---------------------|------------|------------|
| ALEX | 167.383 | 175.213 | 19 | 3206.53 |
| CSS-Btree | 275.09 | 268.488 | 25 | 213.12 |
| DeltaCode16 | 902.359 | 823.844 | 123 | 1464.73 |
| DeltaCode32 | 1034.87 | 966.477 | 114 | 1581.43 |
| EliasFano | 9741.94 | 280.261 | 50 | 1588.23 |
| GammaCode16 | 868.381 | 802.847 | 163 | 1189.7 |
| GammaCode32 | 980.971 | 928.791 | 156 | 1196.68 |
| LA-vector10 | 276.811 | 230.765 | 65 | 530.067 |
| LA-vector12 | 299.576 | 233.954 | 76 | 406.322 |
| LA-vector6 | 399.53 | 431.715 | 114 | 1475.26 |
| LA-vector8 | 314.007 | 332.131 | 69 | 709.681 |
| LA-vectoropt | 365.824 | 276.092 | 54 | 48099.2 |
| PGM++128 | 345.919 | 321.539 | 1 | 84.5145 |
| PGM++32 | 262.665 | 246.796 | 1 | 100.503 |
| PGM++8 | 240.177 | 224.297 | 3 | 136.395 |
| PGM128 | 330.631 | 314.629 | 1 | 85.9683 |
| PGM32 | 258.713 | 240.903 | 1 | 102.395 |
| PGM8 | 237.098 | 221.211 | 3 | 139.217 |
| PLEX128 | 226.382 | 212.227 | 1 | 795.402 |
| PLEX32 | 174.259 | 145.882 | 2 | 903.678 |
| PLEX8 | 147.392 | 126.838 | 18 | 1169.39 |
| RMI-compact | 163.43 | 110.803 | 7 | 1.92038 |
| RMI-large | 110.573 | 102.847 | 403 | 163.962 |
| SIMD-BTree | 107.4 | 88.0679 | 1 | 283.448 |
| SIMD-SampledBTree | 75.8451 | 64.8493 | 13 | 189.147 |
| std::vector | 529.908 | 481.796 | 200 | 0.00030091 |

Table 6: Tabular data: zipf dataset

| index | lookup existing (ns) | lookup missing (ns) | space (MB) | build (ms) |
|-------------------|----------------------|---------------------|------------|-------------|
| ALEX | 101.038 | 89.5688 | 1 | 1566.54 |
| CSS-Btree | 298.307 | 293.062 | 25 | 254.656 |
| DeltaCode16 | 878.739 | 872.093 | 82 | 1372.16 |
| DeltaCode32 | 1002.53 | 999.058 | 72 | 1507.11 |
| EliasFano | 306.291 | 287.017 | 50 | 1552.98 |
| GammaCode16 | 846.473 | 843.143 | 88 | 1054.83 |
| GammaCode32 | 935.681 | 944.724 | 78 | 1078.23 |
| LA-vector10 | 260.203 | 257.764 | 65 | 514.121 |
| LA-vector12 | 271.374 | 260.617 | 76 | 399.656 |
| LA-vector6 | 453.085 | 443.034 | 120 | 1404.77 |
| LA-vector8 | 344.367 | 356.074 | 66 | 633.303 |
| LA-vectoropt | 334.535 | 324.737 | 62 | 48471.2 |
| PGM++128 | 328.339 | 323.115 | 1 | 85.1016 |
| PGM++32 | 258.105 | 254.79 | 1 | 100.473 |
| PGM++8 | 229.672 | 234.455 | 3 | 136.707 |
| PGM128 | 289.289 | 297.185 | 1 | 86.4902 |
| PGM32 | 248.307 | 244.356 | 1 | 99.0724 |
| PGM8 | 234.847 | 229.159 | 3 | 141.955 |
| PLEX128 | 207.873 | 215.56 | 1 | 829.771 |
| PLEX32 | 157.443 | 151.75 | 2 | 953.041 |
| PLEX8 | 126.653 | 131.623 | 18 | 1254.72 |
| RMI-compact | 124.905 | 118.324 | 7 | 1.97878 |
| RMI-large | 117.204 | 120.689 | 403 | 200.492 |
| SIMD-BTree | 107.516 | 103.903 | 1 | 273.753 |
| SIMD-SampledBTree | 76.2583 | 77.1358 | 13 | 179.501 |
| std::vector | 521.889 | 521.671 | 200 | 0.000394881 |

Table 7: Tabular data: normal dataset

| index | lookup existing (ns) | lookup missing (ns) | space (MB) | build (ms) |
|-------------------|----------------------|---------------------|------------|-------------|
| ALEX | 167.491 | 65.8118 | 1 | 8640.33 |
| CSS-Btree | 373.425 | 122.897 | 100 | 895.002 |
| DeltaCode16 | 1111.35 | 594.898 | 182 | 5179.6 |
| DeltaCode32 | 1109.42 | 901.669 | 138 | 5678.04 |
| EliasFano | 485.996 | 263.342 | 173 | 8624.73 |
| GammaCode16 | 1079.56 | 545.488 | 173 | 4141.61 |
| GammaCode32 | 1086.91 | 836.371 | 129 | 4184.85 |
| PGM++128 | 387.687 | 160.025 | 1 | 408.786 |
| PGM++32 | 306.026 | 149.758 | 4 | 538.527 |
| PGM++8 | 262.989 | 136.68 | 23 | 710.456 |
| PGM128 | 363.67 | 145.096 | 1 | 413.025 |
| PGM32 | 297.806 | 130.935 | 4 | 540.04 |
| PGM8 | 262.699 | 124.676 | 23 | 718.454 |
| PLEX128 | 263.166 | 108.248 | 3 | 3356.92 |
| PLEX32 | 191.716 | 102.956 | 20 | 3736.08 |
| PLEX8 | 184.065 | 125.158 | 135 | 4980.63 |
| RMI-compact | 235.018 | 143.763 | 7 | 4.17444 |
| RMI-large | 184.596 | 134.513 | 403 | 316.373 |
| SIMD-BTree | 123.756 | 53.5838 | 1 | 1019.2 |
| SIMD-SampledBTree | 100.765 | 41.9515 | 51 | 666.576 |
| std::vector | 711.094 | 188.601 | 800 | 0.000350364 |

Table 8: Tabular data: amzn uint32 dataset

| index | lookup existing (ns) | lookup missing (ns) | space (MB) | build (ms) |
|-------------------|----------------------|---------------------|------------|------------|
| ALEX | 378.669 | 375.659 | 7 | 51129.6 |
| CSS-Btree | 537.161 | 473.147 | 800 | 13740.9 |
| DeltaCode16 | 1669.93 | 1486.91 | 4554 | 34080.4 |
| DeltaCode32 | 1741.05 | 1641.7 | 4280 | 34163.2 |
| EliasFano | 870.018 | 633.128 | 3689 | 38514.7 |
| GammaCode16 | 1641.88 | 1480.54 | 6529 | 29730.1 |
| GammaCode32 | 1679.7 | 1529.27 | 6320 | 29658.6 |
| LA-vector10 | 1083.72 | 930.558 | 5401 | 88643.9 |
| LA-vector12 | 1062.75 | 852.285 | 5601 | 99877.4 |
| LA-vector6 | 1053.73 | 879.114 | 4801 | 87544.7 |
| LA-vector8 | 1033.01 | 825.049 | 5001 | 88549 |
| PGM++128 | 501.538 | 528.537 | 5 | 2045.27 |
| PGM++32 | 457.759 | 468.919 | 40 | 2536.9 |
| PGM++8 | 463.501 | 449.182 | 273 | 3789.16 |
| PGM128 | 561.801 | 490.928 | 5 | 2093.48 |
| PGM32 | 488.42 | 431.021 | 40 | 2606.65 |
| PGM8 | 437.634 | 429.747 | 273 | 4060.13 |
| PLEX128 | 385.164 | 337.635 | 20 | 15909.1 |
| PLEX32 | 278.458 | 276.138 | 186 | 17986.7 |
| PLEX8 | 233.454 | 229.739 | 775 | 23171.5 |
| RMI-compact | 354.264 | 297.817 | 7 | 2.93933 |
| RMI-large | 210.585 | 193.132 | 403 | 193.07 |
| Roaring | 3334.82 | 2763.85 | 13123 | 240789 |
| SIMD-BTree | 230.589 | 213.628 | 1 | 13592.8 |
| SIMD-SampledBTree | 211.211 | 198.361 | 801 | 13623.6 |
| std::vector | 1105.49 | 1005.29 | 6400 | 0.00035977 |

Table 9: Tabular data: amzn uint64 dataset

| index | lookup existing (ns) | lookup missing (ns) | space (MB) | build (ms) |
|-------------------|----------------------|---------------------|------------|------------|
| ALEX | 671.084 | 237.645 | 116 | 69691.4 |
| CSS-Btree | 490.232 | 234.792 | 800 | 6953.19 |
| DeltaCode16 | 1758.77 | 794.365 | 4007 | 29418.3 |
| DeltaCode32 | 1847.87 | 1004.56 | 3714 | 29451.8 |
| EliasFano | 5752.57 | 1206.97 | 3691 | 40745.9 |
| GammaCode16 | 1749.51 | 771.32 | 5589 | 26382.1 |
| GammaCode32 | 1872.32 | 918.599 | 5349 | 26258 |
| PGM++128 | 767.681 | 361.784 | 11 | 2465.78 |
| PGM++32 | 759.417 | 310.022 | 48 | 2747.44 |
| PGM++8 | 733.479 | 302.344 | 220 | 3698.03 |
| PGM128 | 745.425 | 323.391 | 11 | 2458.18 |
| PGM32 | 863.109 | 338.437 | 48 | 2775.28 |
| PGM8 | 600.297 | 324.233 | 220 | 3746.43 |
| PLEX128 | 501.306 | 216.359 | 34 | 18646 |
| PLEX32 | 510.674 | 175.989 | 147 | 20487.8 |
| PLEX8 | 390.956 | 160.757 | 879 | 26167.9 |
| RMI-compact | 754.173 | 131.459 | 7 | 2.33918 |
| RMI-large | 477.233 | 107.309 | 403 | 194.959 |
| Roaring | 3078.13 | 984.998 | 9099 | 187146 |
| SIMD-BTree | 307.941 | 148.401 | 1 | 8230.16 |
| SIMD-SampledBTree | 216.965 | 109.716 | 801 | 6987.79 |
| std::vector | 1031.15 | 497.749 | 6400 | 0.00029644 |

Table 10: Tabular data: OSM cellids dataset

| index | lookup existing (ns) | lookup missing (ns) | space (MB) | build (ms) |
|-------------------|----------------------|---------------------|------------|-------------|
| ALEX | 166.736 | 146.297 | 7 | 13391.5 |
| CSS-Btree | 403.955 | 236.624 | 200 | 3081.4 |
| DeltaCode16 | 1232.16 | 752.528 | 1113 | 9473.72 |
| DeltaCode32 | 1401.09 | 941.488 | 1095 | 9415.49 |
| EliasFano | 729.419 | 255.153 | 148 | 8688.64 |
| GammaCode16 | 1231.34 | 752.246 | 1782 | 9503.45 |
| GammaCode32 | 1366.65 | 921.002 | 1785 | 9399.67 |
| LA-vector10 | 455.316 | 294.902 | 251 | 3677.02 |
| LA-vector12 | 506.993 | 296.99 | 301 | 3804.49 |
| LA-vector6 | 431.688 | 470.686 | 157 | 3600.5 |
| LA-vector8 | 377.027 | 237.212 | 202 | 3514.79 |
| LA-vectoropt | 527.244 | 309.344 | 111 | 279103 |
| PGM++128 | 438.115 | 282.648 | 1 | 563.241 |
| PGM++32 | 376.922 | 227.561 | 2 | 559.212 |
| PGM++8 | 303.329 | 232.349 | 12 | 645.531 |
| PGM128 | 425.689 | 275.343 | 1 | 583.744 |
| PGM32 | 332.312 | 227.73 | 2 | 575.504 |
| PGM8 | 304.436 | 218.545 | 12 | 674.349 |
| PLEX128 | 281.411 | 195.393 | 2 | 3419.38 |
| PLEX32 | 197.446 | 143.372 | 6 | 3555.53 |
| PLEX8 | 166.412 | 123.024 | 64 | 4248.39 |
| RMI-compact | 192.971 | 111.758 | 7 | 3.77417 |
| RMI-large | 140.415 | 106.22 | 403 | 290.662 |
| SIMD-BTree | 168.845 | 96.0625 | 1 | 2983.48 |
| SIMD-SampledBTree | 139.777 | 83.8916 | 201 | 3248.17 |
| std::vector | 800.183 | 396.749 | 1600 | 0.000284985 |

Table 11: Tabular data: wiki uint64 dataset

| index | lookup existing (ns) | lookup missing (ns) | space (MB) | build (ms) |
|-------------------|----------------------|---------------------|------------|-------------|
| ALEX | 573.6 | 14.5306 | 74 | 18045.2 |
| CSS-Btree | 395.47 | 70.4056 | 200 | 1310.87 |
| DeltaCode16 | 1287.61 | 414.694 | 374 | 6417.91 |
| DeltaCode32 | 1294.78 | 527.001 | 286 | 6509.4 |
| GammaCode16 | 1275.4 | 377.475 | 393 | 5320.56 |
| GammaCode32 | 1259.13 | 461.428 | 305 | 5343.31 |
| PGM++128 | 479.12 | 95.6022 | 5 | 634.277 |
| PGM++32 | 414.727 | 89.9523 | 18 | 730.322 |
| PGM++8 | 396.81 | 113.925 | 71 | 1003.54 |
| PGM128 | 484.534 | 96.599 | 5 | 653.159 |
| PGM32 | 373.327 | 94.9386 | 18 | 753.238 |
| PGM8 | 388.017 | 101.629 | 71 | 1023.94 |
| PLEX128 | 330.364 | 61.1977 | 14 | 4366.79 |
| PLEX32 | 315.18 | 49.0181 | 55 | 4848.92 |
| PLEX8 | 262.922 | 52.3193 | 176 | 6190.32 |
| RMI-compact | 310.753 | 37.2934 | 7 | 2.34779 |
| RMI-large | 213.143 | 37.4332 | 403 | 207.782 |
| Roaring | 676.09 | 44.9812 | 410 | 22295.9 |
| SIMD-BTree | 211.628 | 50.3897 | 1 | 1318.88 |
| SIMD-SampledBTree | 156.053 | 33.8489 | 201 | 1299.44 |
| std::vector | 796.749 | 72.0668 | 1600 | 0.000298861 |

Table 12: Tabular data: FB uint64 dataset

| Compressed Index | Scan 10 | Scan 100 | Scan 1K | Scan 10K |
|------------------|----------|----------|----------|----------|
| DeltaCode16 | 169.902 | 158.382 | 156.902 | 156.262 |
| DeltaCode32 | 290.359 | 266.085 | 262.788 | 262.641 |
| EliasFano | 60.6795 | 56.5687 | 56.0174 | 55.8601 |
| GammaCode16 | 174.914 | 162.492 | 160.746 | 160.167 |
| GammaCode32 | 298.385 | 271.448 | 268.043 | 267.507 |
| LA-vector10 | 47.289 | 43.3315 | 42.4952 | 42.3297 |
| LA-vector12 | 45.0734 | 41.4511 | 40.76 | 40.6897 |
| LA-vector6 | 49.6115 | 46.3757 | 45.3586 | 45.2125 |
| LA-vector8 | 46.0287 | 42.6529 | 41.6661 | 41.5467 |
| LA-vectoropt | 40.5808 | 36.8136 | 35.9605 | 35.7708 |
| std::vector | 0.562015 | 0.220738 | 0.175476 | 0.251958 |

Table 13: Tabular data: SCAN experiments on the companynet dataset (times are expressed in ns)

| Compressed Index | Scan 10 | Scan 100 | Scan 1K | Scan 10K |
|------------------|---------|----------|---------|----------|
| DeltaCode16 | 180.626 | 152.231 | 147.45 | 145.791 |
| DeltaCode32 | 286.332 | 241.294 | 233.723 | 231.971 |
| EliasFano | 67.0721 | 56.6032 | 55.4092 | 55.0058 |
| GammaCode16 | 165.895 | 143.288 | 139.686 | 137.795 |
| GammaCode32 | 272.873 | 229.1 | 223.36 | 221.512 |
| LA-vector10 | 40.9013 | 38.0294 | 37.3074 | 36.7726 |
| LA-vector12 | 42.8927 | 39.0475 | 39.1278 | 38.3544 |
| LA-vector6 | 43.78 | 39.3341 | 38.73 | 38.2743 |
| LA-vector8 | 38.0936 | 34.117 | 34.0408 | 33.5231 |
| LA-vectoropt | 70.3335 | 51.1309 | 49.5432 | 49.0797 |
| std::vector | 2.12109 | 1.38943 | 0.51784 | 0.472924 |

Table 14: Tabular data: SCAN experiments on the wiki uint32 dataset (times are expressed in ns)

| Compressed Index | Scan 10 | Scan 100 | Scan 1K | Scan 10K |
|------------------|---------|----------|----------|----------|
| DeltaCode16 | 82.4027 | 68.7572 | 67.0358 | 66.135 |
| DeltaCode32 | 105.84 | 91.3301 | 89.1847 | 88.4595 |
| EliasFano | 66.3994 | 56.3723 | 55.3097 | 54.8652 |
| GammaCode16 | 82.4171 | 68.9434 | 67.5749 | 66.7945 |
| GammaCode32 | 106.273 | 90.424 | 88.482 | 87.7333 |
| std::vector | 2.05005 | 0.72798 | 0.505714 | 0.435934 |

Table 15: Tabular data: SCAN experiments on the amzn uint32 dataset (times are expressed in ns)

| Compressed Index | Scan 10 | Scan 100 | Scan 1K | Scan 10K |
|------------------|---------|----------|----------|----------|
| DeltaCode16 | 122.292 | 100.978 | 100.201 | 99.5831 |
| DeltaCode32 | 177.565 | 147.963 | 147.04 | 146.48 |
| GammaCode16 | 138.96 | 116.163 | 113.737 | 112.854 |
| GammaCode32 | 203.884 | 168.636 | 165.93 | 165.224 |
| std::vector | 2.35548 | 2.16368 | 0.963471 | 0.925722 |

Table 16: Tabular data: SCAN experiments on the FB uint64 dataset (times are expressed in ns)

Error Report

LA-vectoropt - fb_200M_uint64 - existing: First correction too large
LA-vector6 - fb_200M_uint64 - existing: Bit fields' sizes are not large enough
LA-vector8 - fb_200M_uint64 - existing: Bit fields' sizes are not large enough
LA-vector10 - fb_200M_uint64 - existing: Bit fields' sizes are not large enough
LA-vector12 - fb_200M_uint64 - existing: Bit fields' sizes are not large enough
LA-vectoropt - books_800M_uint64 - existing: First correction too large
LA-vectoropt - osm_cellids_800M_uint64 - existing: First correction too large
LA-vector6 - osm_cellids_800M_uint64 - existing: Segment correction too large
LA-vector8 - osm_cellids_800M_uint64 - existing: Segment correction too large
LA-vector10 - osm_cellids_800M_uint64 - existing: Segment correction too large
LA-vector12 - osm_cellids_800M_uint64 - existing: Segment correction too large
LA-vectoropt - books_200M_uint32 - existing: Bit fields' sizes are not large enough
LA-vector6 - books_200M_uint32 - existing: Bit fields' sizes are not large enough
LA-vector8 - books_200M_uint32 - existing: Bit fields' sizes are not large enough
LA-vector10 - books_200M_uint32 - existing: Bit fields' sizes are not large enough
LA-vector12 - books_200M_uint32 - existing: Bit fields' sizes are not large enough
LA-vectoropt - fb_200M_uint64 - missing: First correction too large
LA-vector6 - fb_200M_uint64 - missing: Bit fields' sizes are not large enough
LA-vector8 - fb_200M_uint64 - missing: Bit fields' sizes are not large enough
LA-vector10 - fb_200M_uint64 - missing: Bit fields' sizes are not large enough
LA-vector12 - fb_200M_uint64 - missing: Bit fields' sizes are not large enough
LA-vectoropt - books_800M_uint64 - missing: First correction too large
LA-vectoropt - osm_cellids_800M_uint64 - missing: First correction too large
LA-vector6 - osm_cellids_800M_uint64 - missing: Segment correction too large
LA-vector8 - osm_cellids_800M_uint64 - missing: Segment correction too large
LA-vector10 - osm_cellids_800M_uint64 - missing: Segment correction too large
LA-vector12 - osm_cellids_800M_uint64 - missing: Segment correction too large
LA-vectoropt - books_200M_uint32 - missing: Bit fields' sizes are not large enough
LA-vector6 - books_200M_uint32 - missing: Bit fields' sizes are not large enough
LA-vector8 - books_200M_uint32 - missing: Bit fields' sizes are not large enough
LA-vector10 - books_200M_uint32 - missing: Bit fields' sizes are not large enough
LA-vector12 - books_200M_uint32 - missing: Bit fields' sizes are not large enough
LA-vectoropt - fb_200M_uint64 - buildtime: First correction too large
LA-vector6 - fb_200M_uint64 - buildtime: Bit fields' sizes are not large enough
LA-vector8 - fb_200M_uint64 - buildtime: Bit fields' sizes are not large enough
LA-vector10 - fb_200M_uint64 - buildtime: Bit fields' sizes are not large enough
LA-vector12 - fb_200M_uint64 - buildtime: Bit fields' sizes are not large enough
LA-vectoropt - books_800M_uint64 - buildtime: First correction too large
LA-vectoropt - osm_cellids_800M_uint64 - buildtime: First correction too large
LA-vector6 - osm_cellids_800M_uint64 - buildtime: Segment correction too large
LA-vector8 - osm_cellids_800M_uint64 - buildtime: Segment correction too large
LA-vector10 - osm_cellids_800M_uint64 - buildtime: Segment correction too large
LA-vector12 - osm_cellids_800M_uint64 - buildtime: Segment correction too large
LA-vectoropt - books_200M_uint32 - buildtime: Bit fields' sizes are not large enough
LA-vector6 - books_200M_uint32 - buildtime: Bit fields' sizes are not large enough
LA-vector8 - books_200M_uint32 - buildtime: Bit fields' sizes are not large enough
LA-vector10 - books_200M_uint32 - buildtime: Bit fields' sizes are not large enough
LA-vector12 - books_200M_uint32 - buildtime: Bit fields' sizes are not large enough
LA-vectoropt - fb_200M_uint64 - scan: First correction too large
LA-vector6 - fb_200M_uint64 - scan: Bit fields' sizes are not large enough

[illegible]